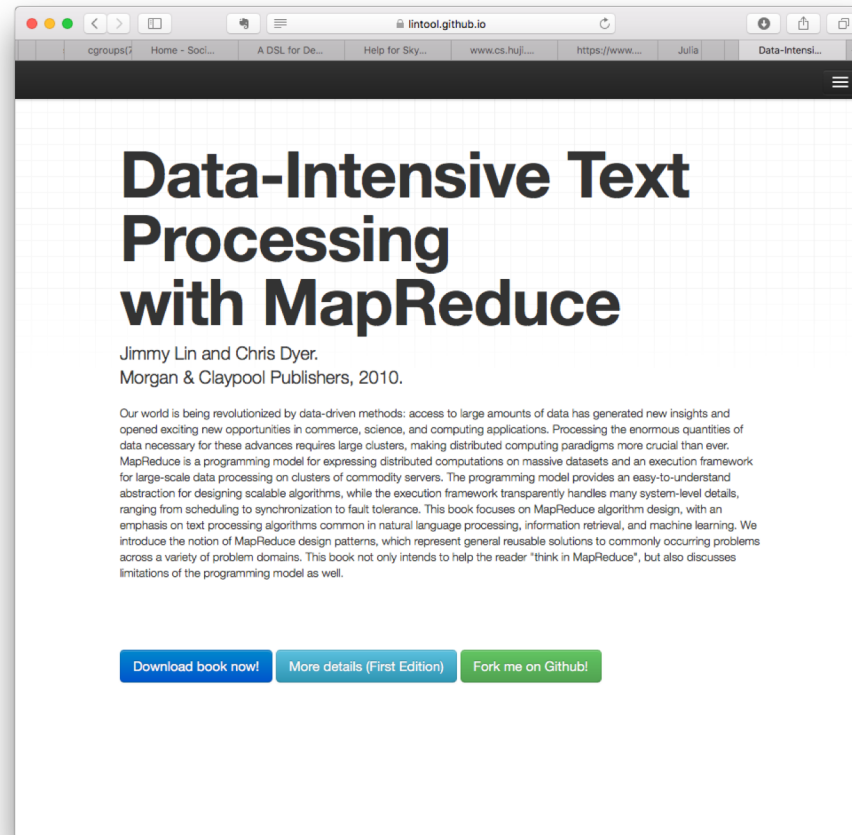


Graphs



Recommended Reading



<https://lintool.github.io/MapReduceAlgorithms/>



Graphs

- Ubiquitous in modern society
 - Hyperlink structure of the Web
 - Social networks
 - Email flow
 - Friend patterns
 - Transportation networks
- Nodes and links can be annotated with metadata
 - Social network nodes: age, gender, interests
 - Social network edges: relationship type and importance

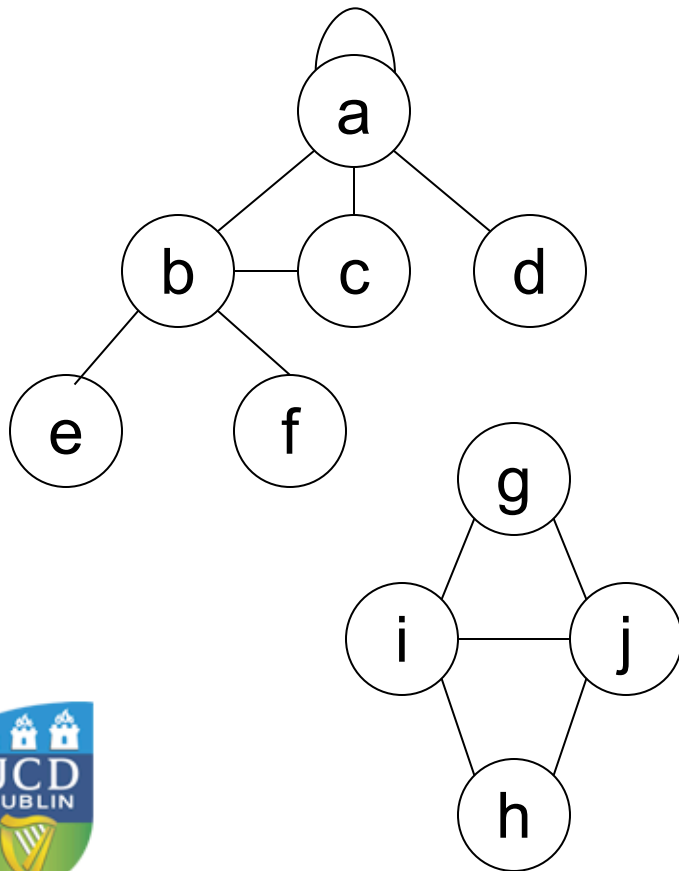


Real-world Problems to Solve

- A common feature: millions or billions of nodes and millions or billions of edges
- Real-world graphs are often sparse, the number of actual edges is far smaller than the number of possible edges

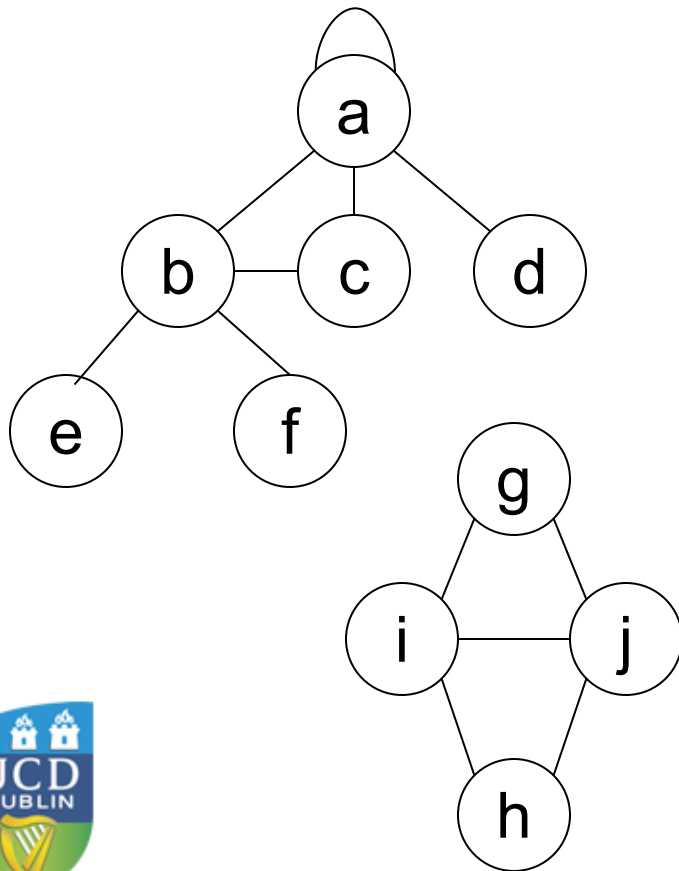


Adjacency List



a	a, b, c, d
b	a, c, e, f
c	a, b
d	a
e	b
f	b
g	i, j
h	i, j
i	g, h, j
j	g, h, i

Adjacency Matrix



	a	b	c	d	e	f	g	h	i	j
a	1	1	1	1						
b	1		1		1	1				
c	1	1								
d	1									
e		1								
f		1								
g									1	1
h									1	1
i							1	1		1
j							1	1	1	

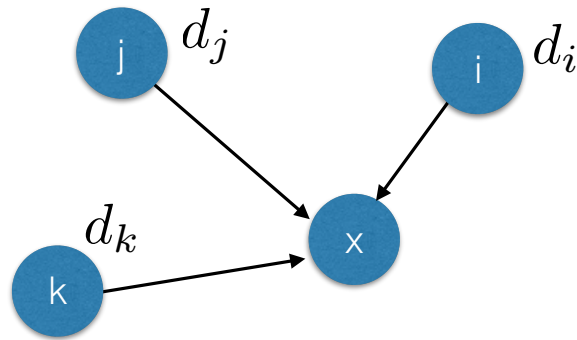
Comparison

- Adjacency Matrix: mathematically easy representation but waste of space
- Adjacency List: A much more compressed representation (for sparse graphs) but some graph operations are more difficult compared to the adj. matrix
- Counting inlinks:
 - Matrix: scan the column and count
 - List: difficult, worst case all data needs to be scanned
- Counting outlinks:
 - Matrix: scan the rows and count
 - List: outlinks are natural



Shortest path in the MapReduce World

- Task: find the shortest path from a source node to all other nodes in the graph. Edges have unit weight.
- Intuition:
 - Distance of nodes N directly connected to the source is 1
 - Distance of nodes directly connected to nodes in N is 2
 - Multiple path to x : the shortest path must go through one of the nodes with an outlink to x ; use the minimum



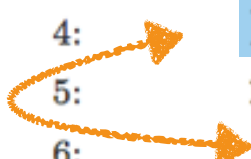
$$d_x = \min(d_i + 1, d_j + 1, d_k + 1)$$



Shortest path in the MapReduce World

```
1: class MAPPER
2:   method MAP(nid  $n$ , node  $N$ )
3:      $d \leftarrow N.DISTANCE$ 
4:     EMIT(nid  $n$ ,  $N$ )
5:     for all nodeid  $m \in N.ADJACENCYLIST$  do
6:       EMIT(nid  $m$ ,  $d + 1$ )

1: class REDUCER
2:   method REDUCE(nid  $m$ , [ $d_1, d_2, \dots$ ])
3:      $d_{min} \leftarrow \infty$ 
4:      $M \leftarrow \emptyset$ 
5:     for all  $d \in \text{counts } [d_1, d_2, \dots]$  do
6:       if ISNODE( $d$ ) then
7:          $M \leftarrow d$ 
8:       else if  $d < d_{min}$  then
9:          $d_{min} \leftarrow d$ 
10:     $M.DISTANCE \leftarrow d_{min}$ 
11:    EMIT(nid  $m$ , node  $M$ )
```



Shortest path in the MapReduce World

- Each iteration of the algorithm is one Hadoop job
 - A map phase to compute the distances
 - A reduce phase to find the current minimum distance
- Iterations:
 1. All nodes connected to the source are discovered
 2. All nodes connected to those discovered in 1. are found
 3. 3. ...
- Between iterations (jobs) the graph structure needs to be passed along; reducer output is input for the next iteration

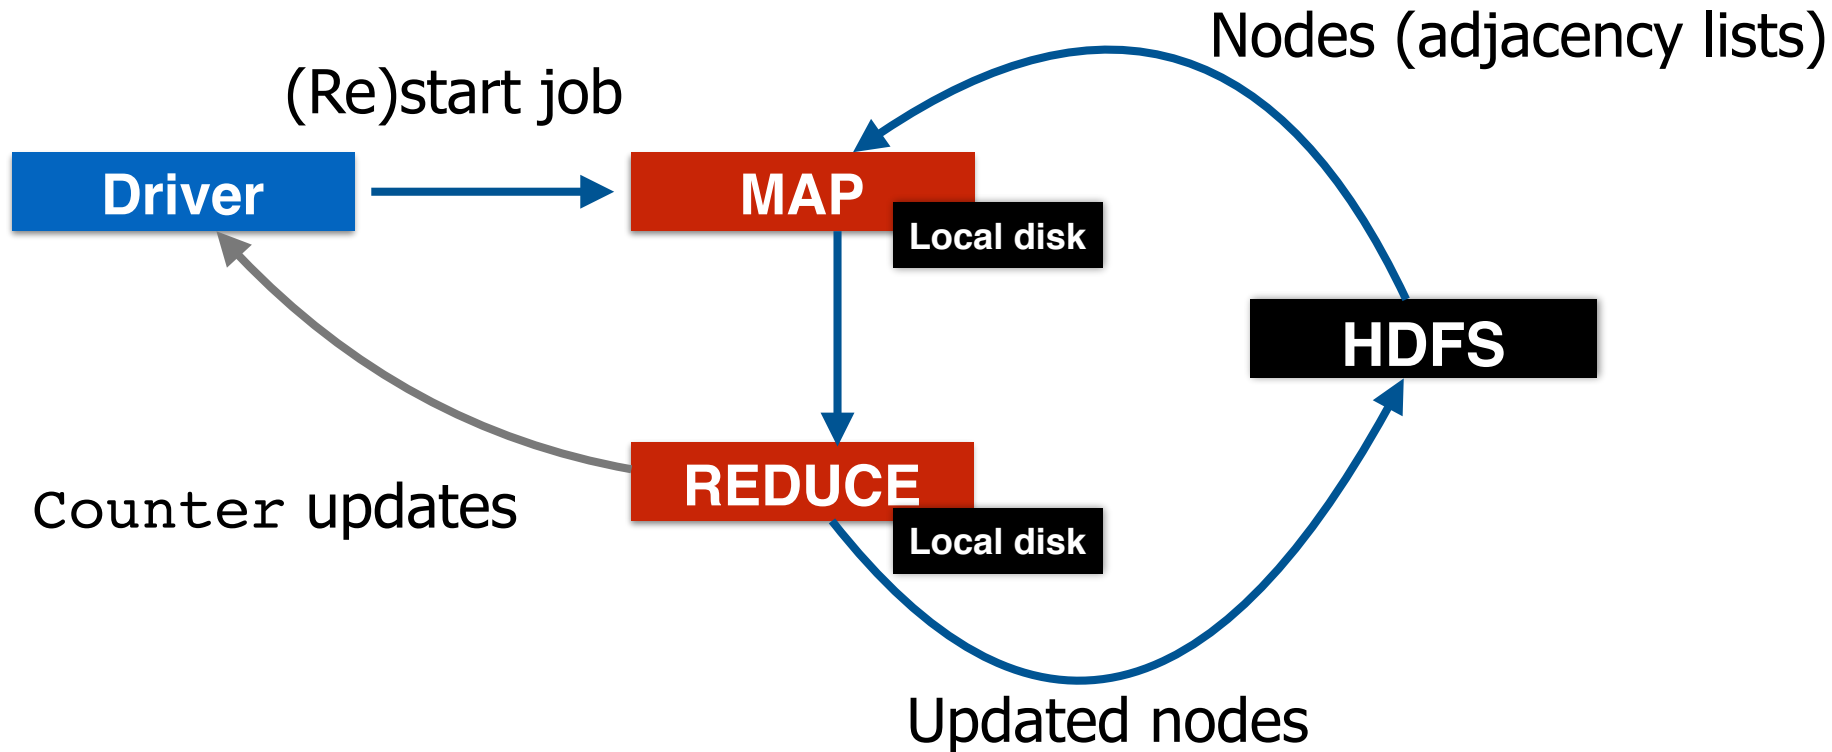


Shortest path in the MapReduce World

- How many iterations are necessary to compute the shortest path to all nodes?
 - Diameter of the graph (greatest distance between a pair of nodes)
 - Diameter is usually small (“six degrees of separation”)
- In practice: iterate until all node distances are less than $+\infty$
 - Assumption: connected graph
- Termination condition checked “outside” of MapReduce job
 - Use Counter to count number of nodes with infinite distance



Shortest path in the MapReduce World



Shortest path in the MapReduce World

- What if the edges have weights?
- Two changes required wrt. the parallel BFS
 - Update rule, instead of $d+1$ use $d+w$
 - Termination criterion: no more distance
- Num. iterations in the worst case: $\#nodes-1$ changes (via Counter)



Single-source Shortest Path

- Dijkstra
 - Single processor (global data structure)
 - Efficient (no recompilation of finalised states)
- Parallel BFS
 - Brute Force approach
 - A lot of unnecessary computations (distances to all nodes recomputed at each iteration)
 - no global data structure



Prototypical approach to graph algorithms in MapReduce/Hadoop

- Node datastructure which contains
 - Adjacency list
 - Additional node [and possibly edge] information (type, features, distances, weights, etc.)
- Job maps over the node data structures
 - Computation involves a node's internal state and local graph structure
 - Result of map phase emitted as values, keyed with node ids of the neighbours; reducer aggregates a node's results
- Graph itself is passed from Mapper to Reducer
- Algorithms are iterative, requiring several Hadoop jobs controlled by the driver code



Outline

- The Web Graph
- PageRank
- Issues and Solutions



THE WEB GRAPH



The Web's Graph Structure

Graph structure in the Web

Andrei Broder^a, Ravi Kumar^{b,*}, Farzin Maghoul^a, Prabhakar Raghavan^b,
Sridhar Rajagopalan^b, Raymie Stata^c, Andrew Tomkins^b, Janet Wiener^c

^a AltaVista Company, San Mateo, CA, USA

^b IBM Almaden Research Center, San Jose, CA, USA

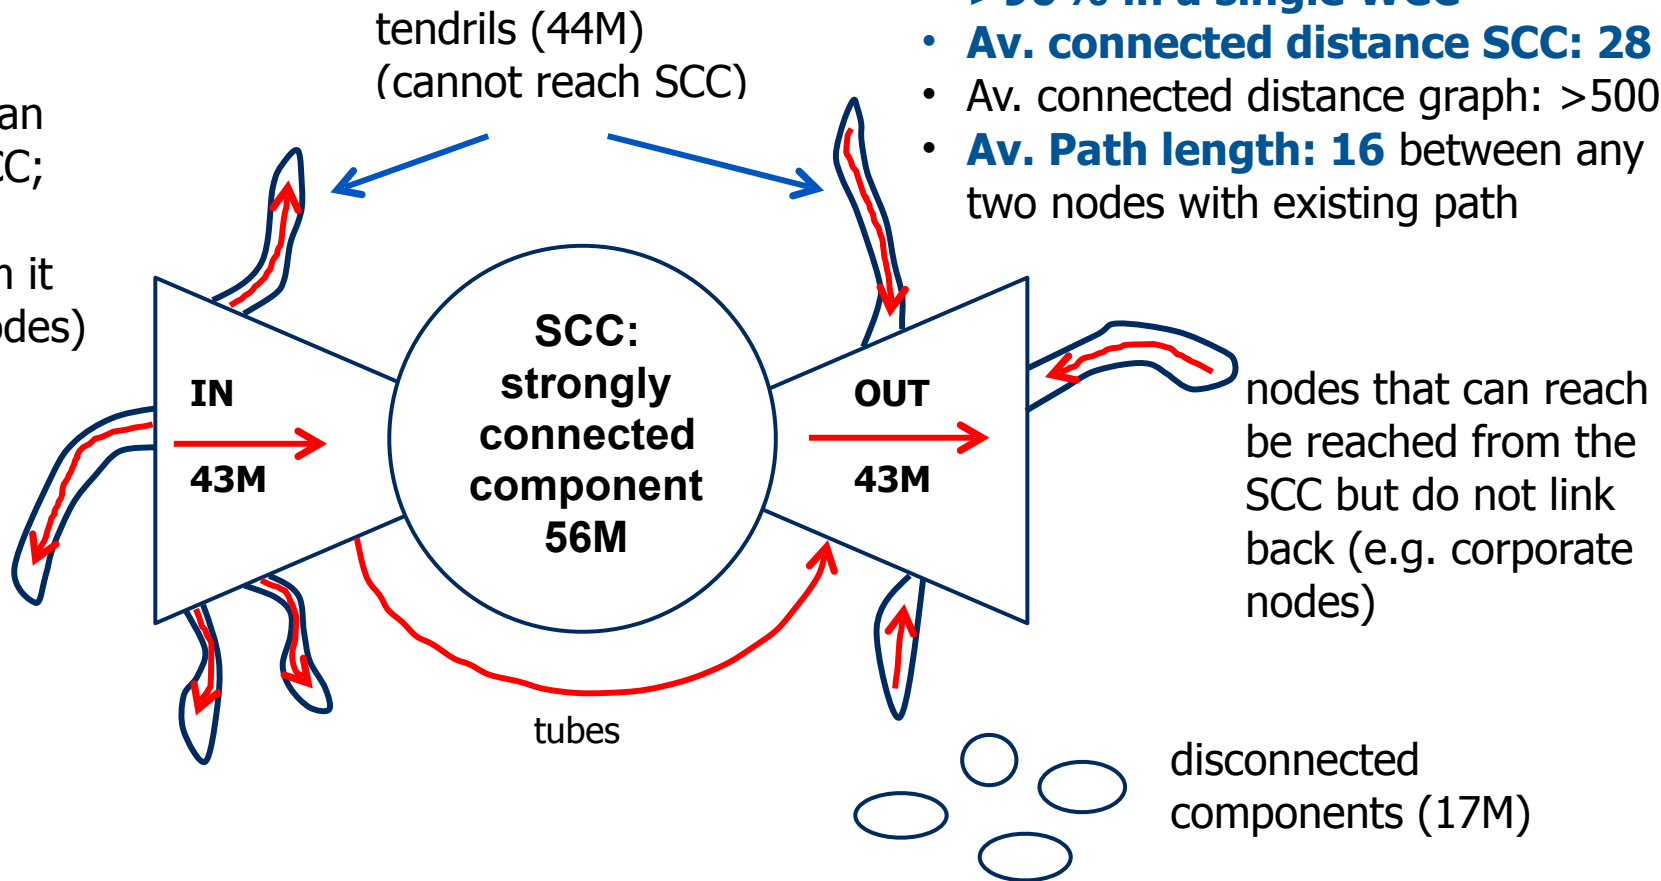
^c Compaq Systems Research Center, Palo Alto, CA, USA

- Insights important for:
 - Crawling strategies
 - Analyzing the behaviour of algorithms that rely on link information (such as PageRank)
 - Predicting the evolution of Web structures
 - etc.
- Data: Altavista crawl from 1999 with 200 million pages and 1.5 billion links



The Web as a “Bow Tie”

nodes that can reach the SCC; cannot be reached from it (e.g. new nodes)



PageRank

- A topic independent approach to page importance
 - Computed once per crawl
- Every document of the corpus is assigned an importance score
 - In search: re-rank (or filter) results with a low PageRank score
- Simple idea: number of in-link indicates importance
 - Page p1 has 10 in-links and one of those is from yahoo.com,
 - page p2 has 50 in-links from obscure pages
- PageRank takes the importance of the page where the link originates into account



“To test the utility of PageRank for search, we built a web search engine called Google.”

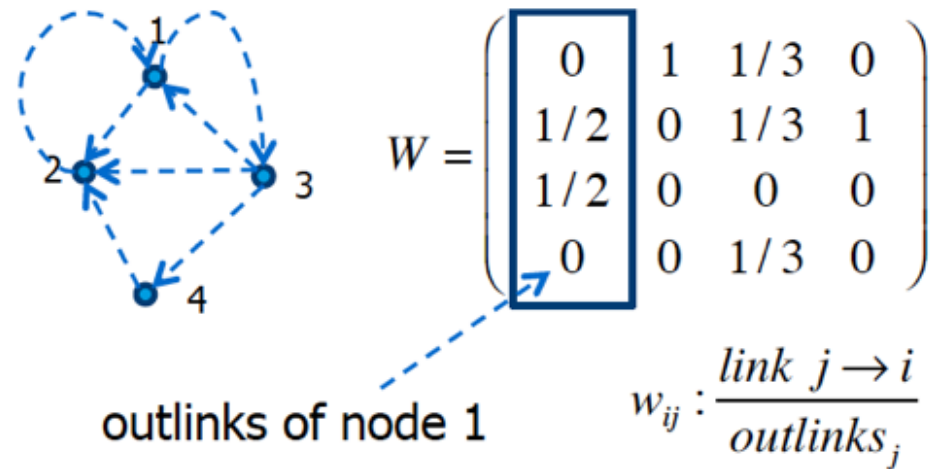
PageRank

- Idea: if page p_x links to page p_y , then the creator of p_x implicitly transfers some importance to page p_y
 - yahoo.com is an important page, many pages point to it
 - Pages linked to from yahoo.com are also likely to be important
- Pages distribute “importance” through outlinks
- Simple PageRank (iteratively)



$$PageRank_{i+1}(v) = \sum_{\substack{u \rightarrow v \\ \text{all nodes linking to } v}} \frac{\text{out-degree of node } u}{PageRank_i(u)} N_u$$

PageRank



Simplified formula

initialize PageRank vector \vec{R}

$$\vec{R} = (R(1), \dots, R(4)) = (0.25, 0.25, 0.25, 0.25)$$

$$W^1 \times \vec{R}' = \begin{pmatrix} 0.33 \\ 0.46 \\ 0.13 \\ 0.08 \end{pmatrix}$$

$$W^2 \times \vec{R}' = \begin{pmatrix} 0.50 \\ 0.29 \\ 0.17 \\ 0.04 \end{pmatrix}$$

$$W^3 \times \vec{R}' = \begin{pmatrix} 0.35 \\ 0.35 \\ 0.25 \\ 0.06 \end{pmatrix}$$

PageRank vector converges eventually

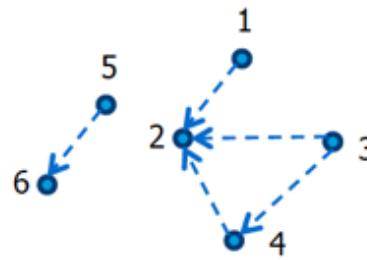
$$W^{16} \times \vec{R}' = \begin{pmatrix} 0.40 \\ 0.33 \\ 0.20 \\ 0.07 \end{pmatrix}$$

$$W^{17} \times \vec{R}' = \begin{pmatrix} 0.40 \\ 0.34 \\ 0.20 \\ 0.07 \end{pmatrix}$$

Random surfer model:

- Probability that a random surfer starts at a random page and ends at page p_x
- A random surfer at a page with 3 outlinks randomly picks one (1/3 prob.) ²²

PageRank



$$W = \begin{pmatrix} 0 & 0 & 0 & 0 \\ 1 & 0 & 1/2 & 1 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 1/2 & 0 \end{pmatrix}$$

Reality

initialize PageRank vector \vec{R}

$$\vec{R} = (R(1), \dots, R(4)) = (0.25, 0.25, 0.25, 0.25)$$

disconnected components

nodes without outgoing edges lead to problems
(rank sink)

$$W^1 \times \vec{R}' = \begin{pmatrix} 0.00 \\ 0.63 \\ 0.00 \\ 0.13 \end{pmatrix}$$

$$W^2 \times \vec{R}' = \begin{pmatrix} 0.00 \\ 0.13 \\ 0.00 \\ 0.00 \end{pmatrix}$$

$$W^3 \times \vec{R}' = \begin{pmatrix} 0.00 \\ 0.00 \\ 0.00 \\ 0.00 \end{pmatrix}$$

Include a decay (“damping”) factor

$$PageRank_{i+1}(v) = \alpha \left(\frac{1}{|G|} \right) + (1 - \alpha) \sum_{u \rightarrow v} \frac{PageRank_i(u)}{N_u}$$

probability that the random surfer
“teleports” and not uses the outlinks

PageRank in MapReduce

- At each iteration:
 - **[MAPPER]** a node passes its PageRank “contributions” to the nodes it is connected to
 - **[REDUCER]** each node sums up all PageRank contributions that have been passed to it and updates its PageRank score

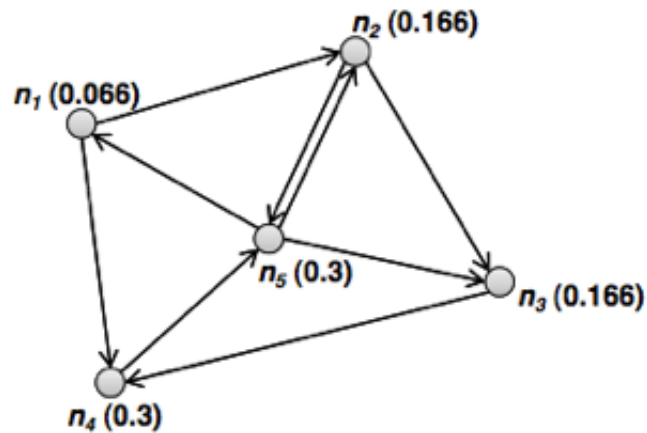
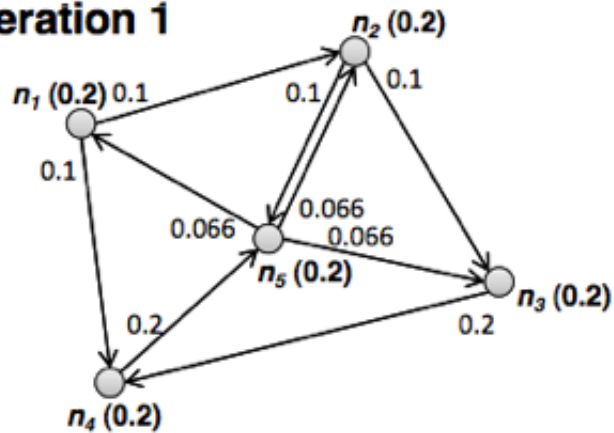


PageRank in MapReduce

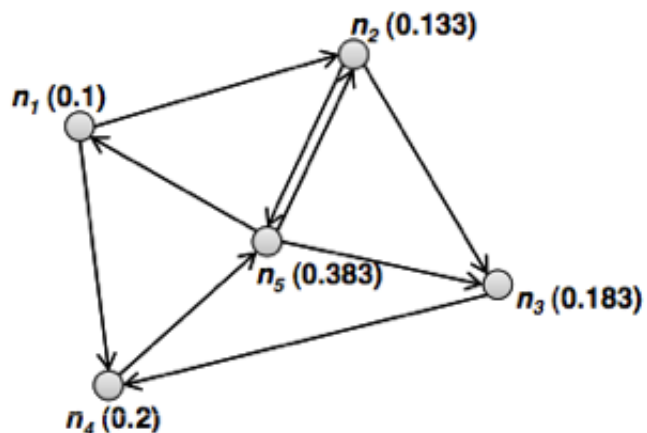
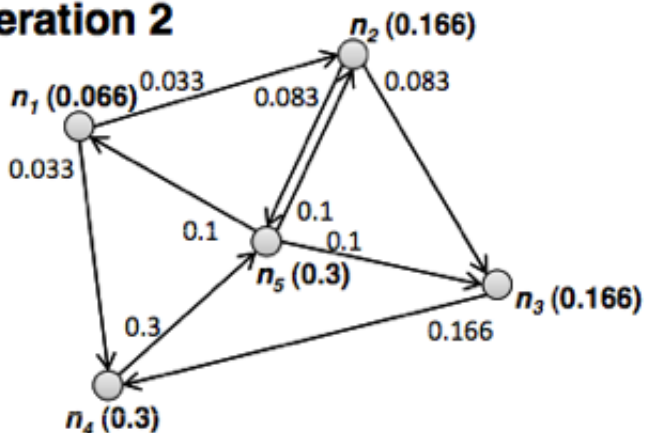
An informal sketch

$$\alpha = 0, \sum_{i=1}^5 n_i = 1$$

Iteration 1



Iteration 2



PageRank in MapReduce

```
1: class MAPPER
2:   method MAP(nid  $n$ , node  $N$ )
3:      $p \leftarrow N.PAGERANK / |N.ADJACENCYLIST|$ 
4:     EMIT(nid  $n$ ,  $N$ )                                ▷ Pass along graph structure
5:     for all nodeid  $m \in N.ADJACENCYLIST$  do
6:       EMIT(nid  $m$ ,  $p$ )                                ▷ Pass PageRank mass to neighbors

1: class REDUCER
2:   method REDUCE(nid  $m$ , [ $p_1, p_2, \dots$ ])
3:      $M \leftarrow \emptyset$ 
4:     for all  $p \in$  counts [ $p_1, p_2, \dots$ ] do
5:       if ISNODE( $p$ ) then
6:          $M \leftarrow p$                                 ▷ Recover graph structure
7:       else
8:          $s \leftarrow s + p$                                 ▷ Sum incoming PageRank contributions
9:      $M.PAGERANK \leftarrow s$ 
10:    EMIT(nid  $m$ , node  $M$ )
```



PageRank in MapReduce

- Dangling nodes: nodes without outgoing edges
 - Simplified PR cannot conserve total PageRank mass (black holes for PR scores)
 - Solution: “lost” PR scores are redistributed evenly across all nodes in the graph
 - Use Counters to keep track of lost mass
 - Reserve a special key for PR mass from dangling nodes
- Redistribution of lost mass and jump factor after each PR iteration in another job (MAP phase only job)



PageRank in MapReduce

- (Possible) stopping criteria
 - PageRank is iterated until convergence (scores at nodes no longer change)
 - PageRank is run for a fixed number of iterations
 - PageRank is run until the ranking of the nodes according to their PR score no longer changes
 - Original PageRank paper: 52 iterations until convergence on a graph with more than 300M edges



ISSUES AND SOLUTIONS



Efficient Large-scale Graph Processing is Challenging

- Poor locality of memory access
- Little work per node (vertex)
- Changing degree of parallelism over the course of execution
- Distribution over many commodity machines due to poor locality is error-prone (failure likely)
- Needed: *“scalable general-purpose system for implementing arbitrary graph algorithms [in batch mode] over arbitrary graph representations in a large-scale distributed environment”*



Existing graph processing options (until 2010)

- Custom distributed infrastructure
 - Problem: each algorithm requires new implementation effort
- Relying on the MapReduce framework
 - Problem: performance and usability issues
 - Remember: the whole graph is read/written in every job
- Single-processor graph algorithm library (e.g. LEDA)
 - Problem: does not scale
- Existing parallel graph systems
 - Problem: do not address fault tolerance & related issues appearing in large distributed setups

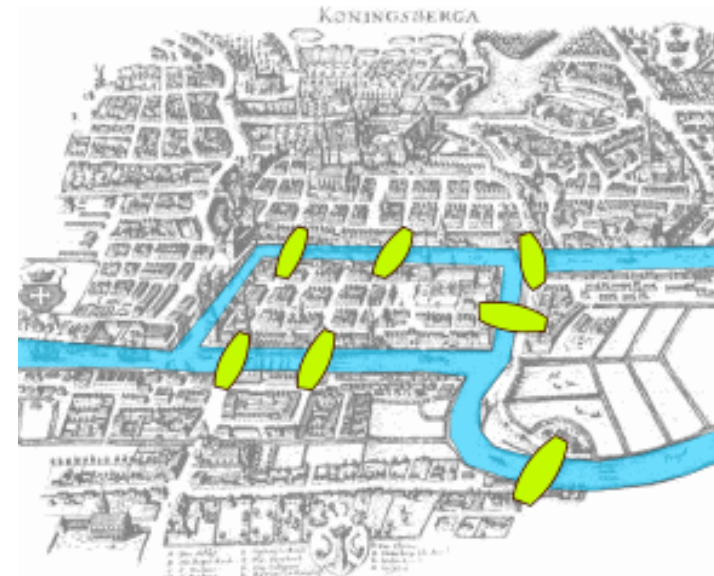


Enter Pregel (2010)

Pregel: A System for Large-Scale Graph Processing

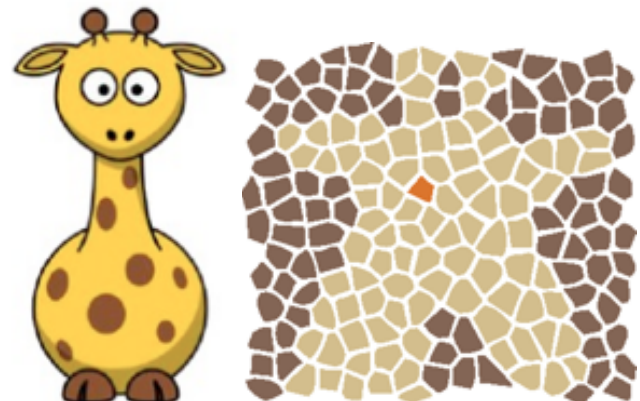
Grzegorz Malewicz, Matthew H. Austern, Aart J. C. Bik, James C. Dehnert, Ilan Horn,
Naty Leiser, and Grzegorz Czajkowski
Google, Inc.
{malewicz,austern,ajcbik,dehnert,ilan,naty,gczaj}@google.com

- “We built a scalable and fault-tolerant platform with an API that is sufficiently flexible to express arbitrary graph algorithms”
- Pregel river runs through Königsberg (Euler’s seven bridges problem)



Graph processing in Hadoop

- Disadvantage: iterative algorithms are slow
 - Lots of reading/writing to and from disk
- Advantage: no additional libraries needed
- Enter Giraph: an open-source implementation of yet another Google framework (Pregel)
 - Specifically created for iterative graph computations



Graph processing in Hadoop

- “Many distributed graph computing systems have been proposed to conduct all kinds of data processing and data analytics in massive graphs, including Pregel, Giraph, GraphLab, PowerGraph, GraphX, Mizan, GPS, Giraph++, Pregelix, Pregel+, and Blogel.”

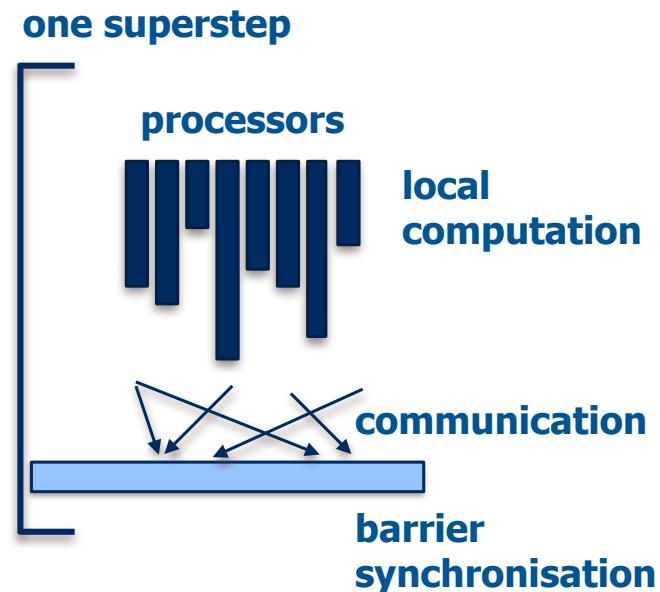


BULK SYNCHRONOUS PARALLEL -- BSP



Bulk Synchronous Parallel

- General model for the design of parallel algorithms
- Developed by Leslie Valiant 1980s/90s
- BSP computer: processors with fast local memory are connected by a communication network
- BSP computation = series of supersteps



- No message passing in MR
- Avoids MR's costly disk and network operations

Bulk Synchronous Parallel

- Supersteps consist of three phases
 1. Local computation: every processor performs computations using data stored in local memory - independent of what happens at other processors; a processor can contain several processes (threads)
 2. Communication: exchange of data between processes (put and get); one-sided communication
 3. Barrier synchronisation: all processes wait until everyone has finished the communication step
- Local computation and communication phases are not strictly ordered in time



Bulk Synchronous Parallel

- BSP & graphs: “**Think like a vertex!**”
- In BSP, algorithms are implemented from the viewpoint of a **single vertex** in the input graph performing a **single iteration** of the computation.

