



NETWORK PROGRAMMING: ACTOR PROGRAMMING

COMP 30220: Distributed Systems

Lecturer: Rem Collier

Email: rem.collier@ucd.ie

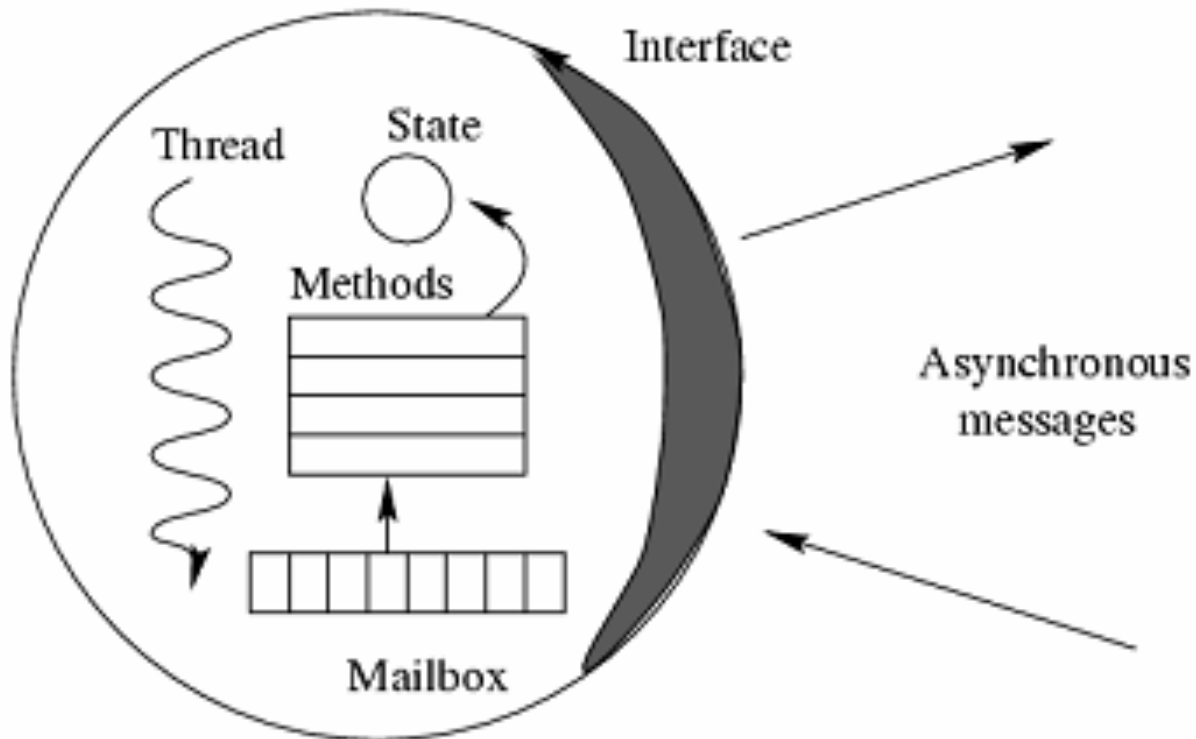
ACTORS

- Initially Proposed in the mid 1970s as a alternative model of concurrent computation.
 - Systems are made up of actor components that interact by passing messages.
 - Computation is realised through the handling of messages by:
 - Making local decisions
 - Creating more actors
 - Sending messages to other actors
 - Actors communicate **asynchronously**:
 - An actor does not wait for a response when sending a message.
 - Actors are **stateful** and **isolated**:
 - They maintain local (private) state
 - They share state by sending messages (no shared memory)
 - Actors are **location independent**:
 - Names do not reflect their physical location.



ANATOMY OF AN ACTOR

- Actor = state + behaviour + mailbox + autonomy



ACTOR SUPERVISION

- Every actor created has a **supervisor** - the actor that created it
 - This means that actor systems are organised hierarchically.
 - Apart from the root actor, every actor has exactly 1 supervisor and 0 or more children.
- Supervisors create actors to perform tasks either on their behalf (delegation) or in response to a larger system need.
- Supervisors play a key role in failure handling:
 - When an actor fails, it suspends all of its children and notifies its supervisor of the failure.
 - The supervisor chooses how to respond to the failure by either resuming the actor (and its children) or restarting the actor.
 - When restarted, the complete actor hierarchy rooted in the failed actor is destroyed.



ACTOR LANGUAGES / LIBRARIES

- Erlang (Telephone exchange programming)
- RevActor (Ruby)
- SALSA
- Kilim
- Reactors
- Akka (Scala and Java)
- Project Orleans (Microsoft Virtual Actors)
- ...





The left side of the slide features a series of vertical stripes in shades of brown, tan, and grey. Overlaid on these stripes are several orange circles of varying sizes, arranged in a cluster that tapers towards the bottom.

THE REACTIVE MANIFESTO

WHY CONSIDER ACTORS?

- 3-tier architecture proven solution to web applications:
 - Data Layer: Code to access the database
 - Service Layer: Main business functions
 - Web Layer: Front end (views and controllers)
- Scaling can be achieved by replicating the 3 layers as necessary.
 - But this does not work for the DB...
 - Can be scaled somewhat using master-slave, clustering, ...
 - Ultimately there is still a single master that must be updated sequentially...



THE TRADITIONAL APPROACH IS NO LONGER ADEQUATE FOR MODERN SYSTEMS

THE REACTIVE MANIFESTO

([HTTP://WWW.REACTIVEMANIFESTO.ORG/](http://www.reactivemanifesto.org/))

- Published 16th September, 2014 (> 20,000 signatures)
- Current approaches to building large-scale software systems is inadequate for the task at hand.
 - The last 10 years has seen a move from systems containing 10s of servers handling gigabytes of data to cloud-based systems that expand to consume 1,000s of servers handling petabytes of data.
- Increased scale makes manual failure monitoring and recovery impractical.
 - Need to develop systems that are designed to automatically detect and recover from component failures.
- The profusion of mobile devices and global nature of many web-based services has led to widely varying usage patterns that exhibit rapid changes in demand for services.
 - Systems need to be designed that can handle wild fluctuations in workload by (re-)allocating resources dynamically and seamlessly



THE REACTIVE MANIFESTO

([HTTP://WWW.REACTIVEMANIFESTO.ORG/](http://www.reactivemanifesto.org/))

- The solution is to promote architectures that are:
 - **Responsive:** Systems should respond in a timely and consistent manner; where possible, providing guarantees on response time.
 - **Resilient:** Failure should be expected and handled transparently by systems while maintaining expected levels of responsiveness.
 - **Elastic:** Systems should remain responsive even as workload varies through dynamic allocation of resource as required.
 - **Message-driven:** Asynchronous message passing promotes loosely-coupled systems. Use of message-passing promotes load-management, elasticity, and flow control through monitoring and management of message queues.



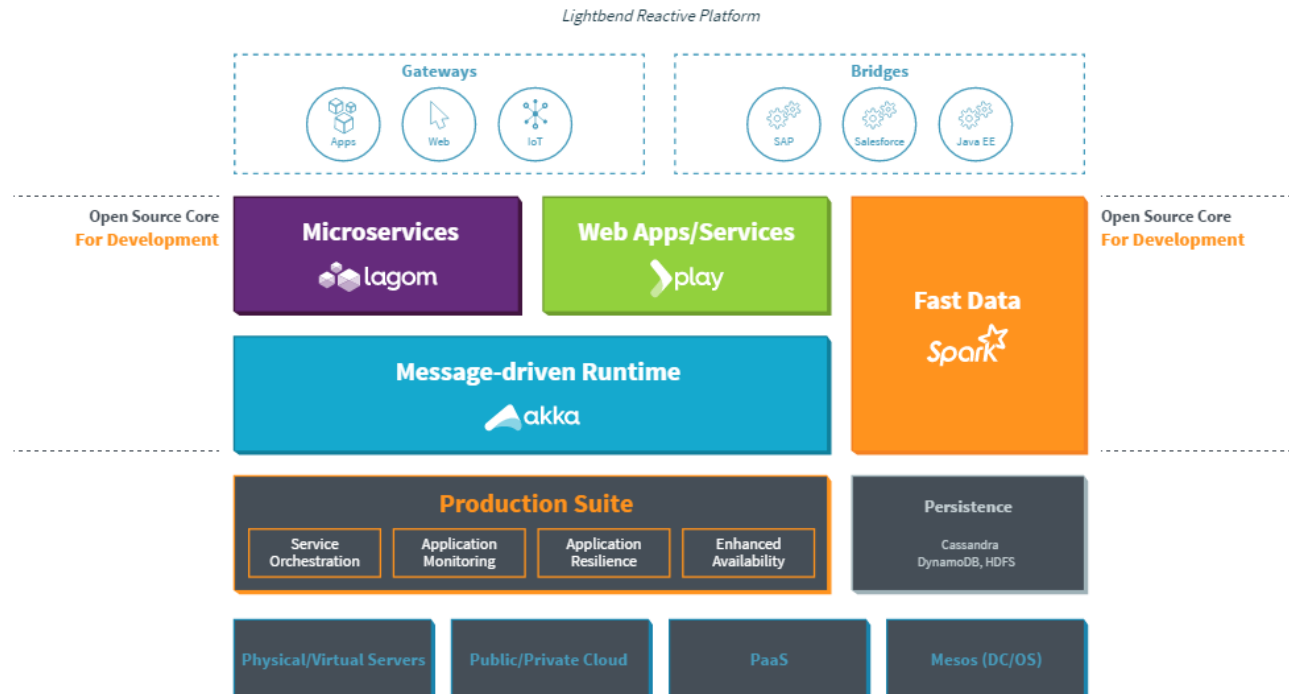
REACTIVE (NEW WEB) TECHNOLOGIES

- Non-blocking IO (NIO):
 - Remove the need to block while waiting to read/write through the use of callbacks / event-queues (streams).
 - *Example instead of blocking, you insert a retry event set for some point in the future. This allows the system to continue to process other events in the meantime....*
- Reactive Programming (Rx)
 - Programming paradigm based on data flow (stream) processing.
 - Handlers perform operations on streams, such as consuming, merging, aggregating, filtering, creating, ...
- Example: Netty (<http://netty.io>)
 - Non-blocking IO/Reactive Programming implementation for Java




EXAMPLE: LIGHTBEND FRAMEWORK

- Typesafe's “New Web” Reactive Framework



<https://www.lightbend.com/platform/development>



INTRODUCING AKKA

AKKA

([HTTP://AKKA.IO/](http://akka.io/))

- A toolkit and runtime for Actor Systems
 - Supported in both Java, Scala, and .NET.
 - Provided as a core library of Scala.
 - Part of the Lightbend framework.
 - Simple to use...
- Widely used in industry:
 - Some Case Studies:
<http://www.lightbend.com/resources/case-studies-and-stories>
 - A good presentation on Akka and “the web”:
<https://www.infoq.com/presentations/Akka-Actors>
- Freely available and open-source...



AKKA CONCEPTS

- To create an actor in Java, you simply extend the `akka.actor.UntypedActor` class and implement the `onReceive(...)` method:

```
public class Hello extends UntypedActor {  
    @Override  
    public void onReceive(Object message) {  
    }  
}
```

- The `onReceive(...)` method is invoked whenever there is an unhandled message in the actors message queue and the actor is idle (not processing another message).
- It is important to note that messages are Java objects...



HELLO AKKA

- To create a “hello world” actor system, we need to do a couple of things:
 1. Create a Hello message that will be sent to the Hello actor:

```
public class HelloMessage {  
}
```

2. Write some code in the `onReceive(...)` method to handle the message and print out the text:

```
public class Hello extends UntypedActor {  
    @Override  
    public void onReceive(Object message) {  
        if (message instanceof HelloMessage) {  
            System.out.println("Hello World");  
        }  
    }  
}
```

3. Write some code to deploy our actor and send the message...



HELLO AKKA

3. Write some code to deploy our actor and send the message...

```
public class Main {  
    public static void main(String[] args) {  
        ActorSystem system =  
            ActorSystem.create("ContentSystem");  
  
        final ActorRef hello =  
            system.actorOf(  
                Props.create(Hello.class),  
                "hello");  
  
        hello.tell(new HelloMessage(), null);  
    }  
}
```



HELLO AKKA

3. Write some code to deploy our actor and send the message...

```
public class Main {  
    public static void main(String[] args) {  
        ActorSystem system =  
            ActorSystem.create("ContentSystem");  
  
        final ActorRef hello =  
            system.actorOf(  
                Props.create(Hello.class),  
                "hello");  
  
        hello.tell(new HelloMessage(), null);  
    }  
}
```

1. Create the Actor System

2. Create the "Hello" Actor

3. Send the initial message



ASIDE: SETTING UP AN AKKA PROJECT

- Download the Akka distribution from <http://akka.io>
- Create a new Java Project
- Add the following jar files from the akka distribution to your classpath (XXX does not need to be the same for each file):
 - akka-actor-XXX.jar
 - scala-library-XXX.jar
 - config-XXX.jar
- You are ready to go ☺



A BIGGER EXAMPLE

- Problem: Concurrent Summation of Values
 - Take a list of integers, break it into sub-lists that are distributed to actors who return the sum of the values in each sub-list.
- Solution:
 - Two actor types: Master (problem creator) and Worker (problem solver)
 - Master will receive an Init message defining:
 - How many random numbers
 - How many workers
 - Master will create the array of numbers and the workers
 - Master breaks the list into segments and passes one segment to each worker using a Problem message
 - Workers sum the values in the segment and return the result using a Result message.
 - Master combines the answers to give the final total



A BIGGER EXAMPLE: MESSAGES

- Init Message:

```
public class Init {  
    public int numbers;  
    public int workers;  
  
    public Init(int numbers, int workers) {  
        this.numbers = numbers;  
        this.workers = workers;  
    }  
}
```

- Problem Message:

```
public class Problem {  
    public int[] list;  
    public int start;  
    public int end;  
  
    public Problem(int[] list, int start, int end) {  
        this.list = list;  
        this.start = start;  
        this.end = end;  
    }  
}
```



A BIGGER EXAMPLE: MESSAGES

- Result Message:

```
public class Result {  
    public int value;  
  
    public Result(int value) {  
        this.value = value;  
    }  
}
```



A BIGGER EXAMPLE: WORKER

```
public class Worker extends UntypedActor {  
    @Override  
    public void onReceive(Object message) throws Throwable {  
        if (message instanceof Problem) {  
            Problem problem = (Problem) message;  
  
            // Do the calculation  
            int sum = 0;  
            for (int i=problem.start; i<problem.end; i++) {  
                sum += problem.list[i];  
            }  
  
            // Send the result  
            getSender().tell(new Result(sum), getSelf());  
            getSelf().tell(Kill.getInstance(), getSelf());  
        }  
    }  
}
```



A BIGGER EXAMPLE: MASTER

```
public class Master extends UntypedActor {  
    Random random = new Random();  
    ActorRef[] workers = null;  
    int count;  
    int sum;  
    long startTime;  
  
    @Override  
    public void onReceive(Object message) throws Throwable {  
        if (message instanceof Init) {  
            Init init = (Init) message;  
  
            // Create the workers  
            workers = new ActorRef[init.workers];  
            for (int i=0; i<init.workers; i++) {  
                workers[i] = getContext().actorOf(  
                    Props.create(Worker.class), "worker_"+i);  
            }  
        }  
    }  
}
```



A BIGGER EXAMPLE: MASTER

```
// Create the list
int[] list = new int[init.numbers];
for (int i=0; i<init.numbers;i++) {
    list[i]=random.nextInt(100);
}

// Local Summation
int lsum = 0;
long st = System.currentTimeMillis();
for (int i=0; i<init.numbers;i++) {
    lsum+= list[i];
}
System.out.println("local result = " + lsum);
System.out.println("local time = " +
    (System.currentTimeMillis()-st) + "ms");

// Start the actor based approach
startTime = System.currentTimeMillis();
```



A BIGGER EXAMPLE: MASTER

```
// Distribute the problems
int segment = init.numbers / init.workers;
for (int i=0; i<init.workers; i++) {
    workers[i].tell(
        new Problem(list, i*segment, ((i+1)*segment)-1),
        getSelf());
}
count = 0; sum = 0;
} else if (message instanceof Result) {
    sum += ((Result) message).value; count++;
    if (count == workers.length) {
        System.out.println("result = " + sum);
        System.out.println("duration = " +
            (System.currentTimeMillis()-startTime) + "ms");
        getSelf().tell(Kill.getInstance(), getSelf());
    }
}
}
}
```



WHAT ABOUT DISTRIBUTION?

- Akka supports a number of distribution models:
 - Remote: Linking of two akka runtimes based on TCP/IP.
 - Permits remote referencing of actor using ActorSelection class.
 - Permits remote creation of actors using Deploy class.
 - Apache Camel: Integration framework (supports SOAP, HTTP, JMS, ...)
 - Akka integration allows you to easily create and deploy distributed actor systems.
 - Clustering: Decentralised peer-to-peer membership service.
 - Uses a gossip protocol to maintain membership information
 - Vector clocks for event ordering
 - Distributed failure monitoring used to monitor for failed runtimes

