



DISTRIBUTED OBJECTS

COMP 30220: Distributed Systems

Lecturer: Rem Collier

Email: rem.collier@ucd.ie

INTRODUCTION

- **Idea:** Distributed Objects are an Object-Oriented approach to implementing Remote Procedure Calls.
 - A system consists of a set of distributed objects.
 - Objects are identified by a unique identifier.
 - Computation involves the invocation of methods on specified objects.
 - This is done transparently whether the object is local or remote.
 - Distributed Objects Frameworks try to hide the network from the programmer.
- Distributed Object Frameworks:
 - CORBA – Common Object Request Broker Architecture
 - DCOM – Distributed Common Object Model
 - Java RMI – Java Remote Method Invocation
 - JINI – Apache Distributed Objects Project (Now Apache River)





COMMON OBJECT REQUEST BROKER ARCHITECTURE (CORBA)

INTRODUCTION TO CORBA

- CORBA is a standard for enabling interaction between applications written in heterogenous languages.
 - Based on the definition of public interfaces written the Interface Definition Language (IDL).

```
module Finance {  
    typedef sequence<string> StringSeq;  
    struct AccountDetails {  
        string name;  
        StringSeq address;  
        long account_number;  
        double current_balance;  
    };  
    exception insufficientFunds { };  
    interface Account {  
        void deposit(in double amount);  
        void withdraw(in double amount) raises(insufficientFunds);  
        readonly attribute AccountDetails details;  
    };  
};
```

- Language specific calls are mapped to the IDL interface:
 - Mappings exists for: C, C++, Ada, Cobol, Java, Lisp, Ruby, Python,
...



IDL MAPPINGS

- Primitive types are mapped to their equivalent in the target language
 - This includes viewing `string` as a primitive type
- Other more complex types include:
 - `struct` types are mapped to C structs or classes containing only public fields.
 - `sequence` defines a vector type.
 - `typedef` is used to define new names for existing types
 - `union` defines a type that can hold several values at runtime
 - In Java, this is a class containing 1 field per type and a mechanism to check which type is returned – discriminator).
 - `interface` defines the distributed objects exposed by CORBA.
 - `module` defines a set of objects (like a package/namespace)




USING CORBA

- Application development starts with an IDL definition.
- An IDL Compiler is used to generate code in a specific (set of) target language(s).
 - structs and unions are mapped to Java classes
 - Field types are inferred from primitive types, typedefs, enums and sequences.
 - The interface is mapped to a triplet of:
 - A **Java interface** of the same name
 - **POA(Portable Object Adaptor)/Skeleton code** that implements the server side of the network interface (and which is extended for specific implementations).
 - **Stub/Proxy code** that implements the client side (may not be needed).
 - All classes are organised within packages based on module names.




EXECUTING A CORBA RPC



- The Server creates deploys a CORBA object by:
 1. creating an ORB,
 2. creating a POA,
 3. creating the servant object,
 4. registering the POA/servant with the ORB, and
 5. advertising the POA/servant via the **COSNaming** service.



EXECUTING A CORBA RPC



- The Client uses a CORBA object by:
 1. locating the reference of the remote object via the COSNaming Service
 2. retrieving the proxy that represents the remote object
 3. Invoking a method on the proxy that causes the equivalent method on the remote object to be invoked.



IIOP: INTERNET INTER-ORB PROTOCOL

- IIOP = GIOP + TCP/IP
- GIOP (General Inter-ORB Protocol) is an abstract protocol:
 - Defines byte-based data transfer syntax called Common Data Representation (CDR).
 - Supports 8 basic message types:
 - Request, Reply, CancelRequest, LocateRequest, LocateReply, CloseConnection, MessageError, Fragment.
- IIOP does not support Fragment or MessageError
- IIOP is now used beyond the scope of CORBA



THE RISE AND FALL OF CORBA

- CORBA pre-dates the rise of the Web.
 - In the 1990s, it was a leading technology for the creation of distributed applications.
- Its large suite of standards made the creation of fully standards compliant implementations difficult.
 - Full implementations were required in **each** target language.
 - Many implementations were considered poor.
 - Good implementations were licensed rather than open-sourced.
- Compared to newer technologies, CORBA was found to be lacking:
 - Code was not transferrable between ORBs
 - IIOP required exceptions to be added to firewalls (in contrast with SOAP)
 - Many projects did not require multi-language support and SOA/SOAP seemed easier for those that did.
 - CORBA-based systems often viewed as too brittle for Internet-scale applications.

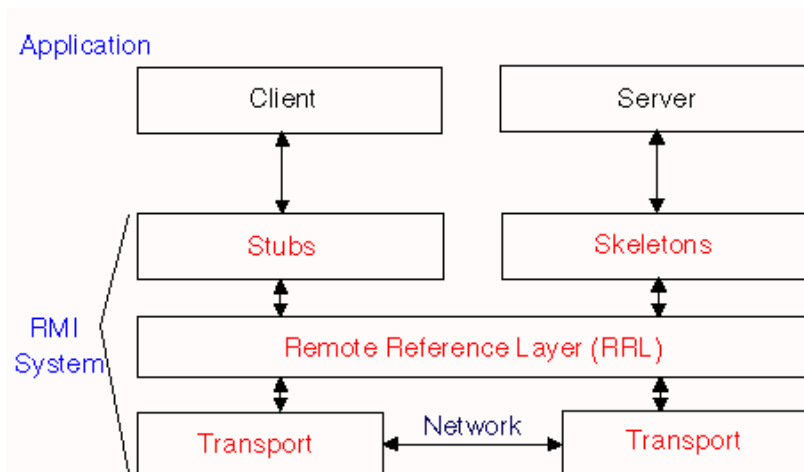


A decorative graphic on the left side of the slide. It features a series of vertical stripes in shades of brown, tan, and grey. Overlaid on these stripes are several orange circles of varying sizes, arranged in a cluster that tapers downwards.

REMOTE METHOD INVOCATION

REMOTE METHOD INVOCATION (RMI)

- Idea: A mechanism to allow Java objects in different JVMs to communicate “directly” with each other
 - A remote object lives on a server and has a remote interface that specifies which of its methods can be invoked by clients.
 - Local objects interact with the remote object by invoking methods on a local “stub” that deals with invoking the method on the remote object.
 - The stub and skeleton (the server side part of the remote connection) are created automatically.



INSERT: SERIALISATION

- Mechanism for converting an object into a stream of bytes:
 - Non-trivial as objects may be composed of other objects...
 - The stream of bytes can be sent anywhere: a file, a server, ...
 - Once received, a copy of the original object can be recreated from the stream of bytes.
 - Actually, the destination device will also need the SAME class definition.
- Java supports serialization of primitive types and objects that implement the `java.io.Serializable` interface.
 - E.g. String, Vector, Integer, Float, Exceptions, most of AWT and Swing...
 - For full list, consult the Javadoc:
<http://docs.oracle.com/javase/9/docs/api/java/io/Serializable.html>

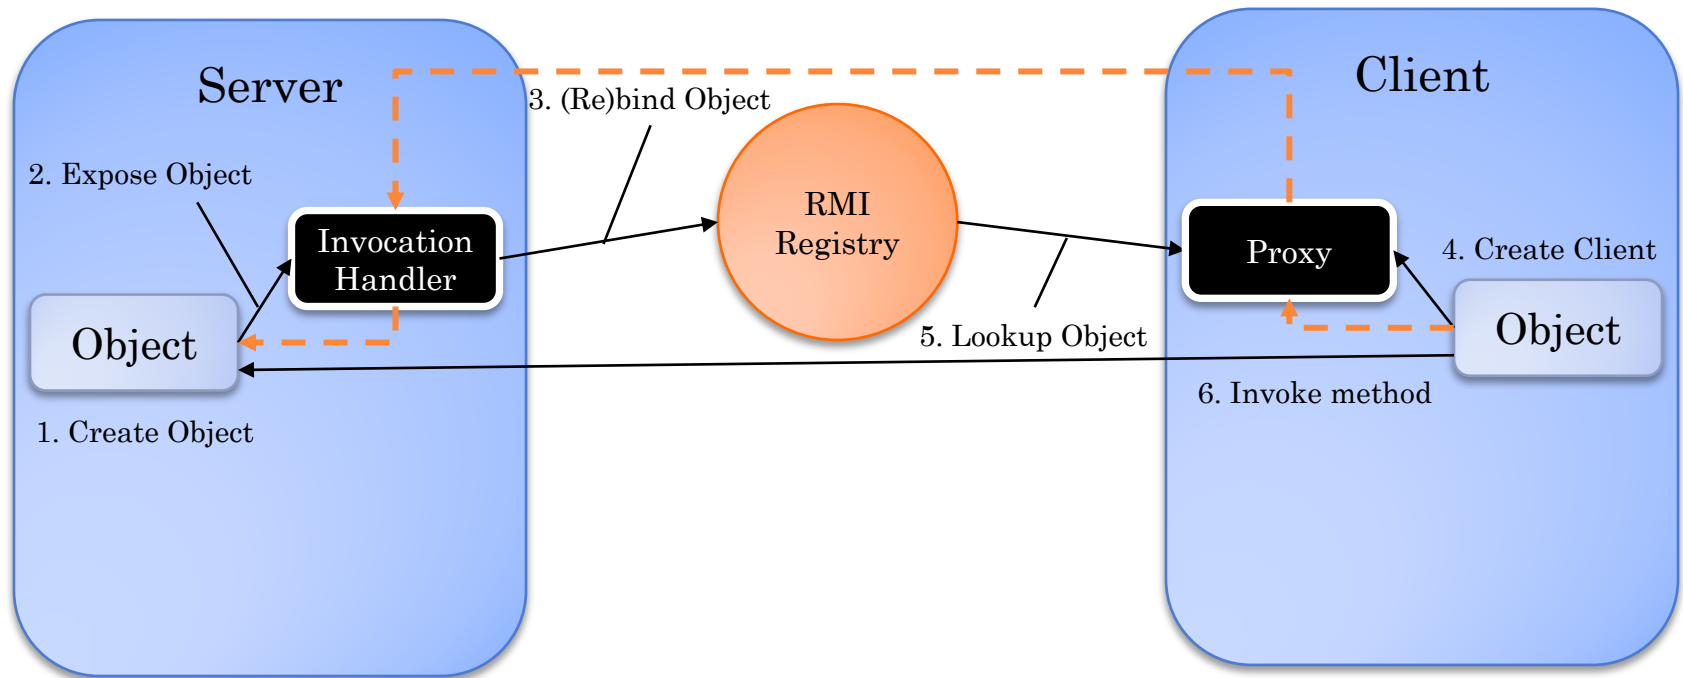


JAVA RMI ARCHITECTURE

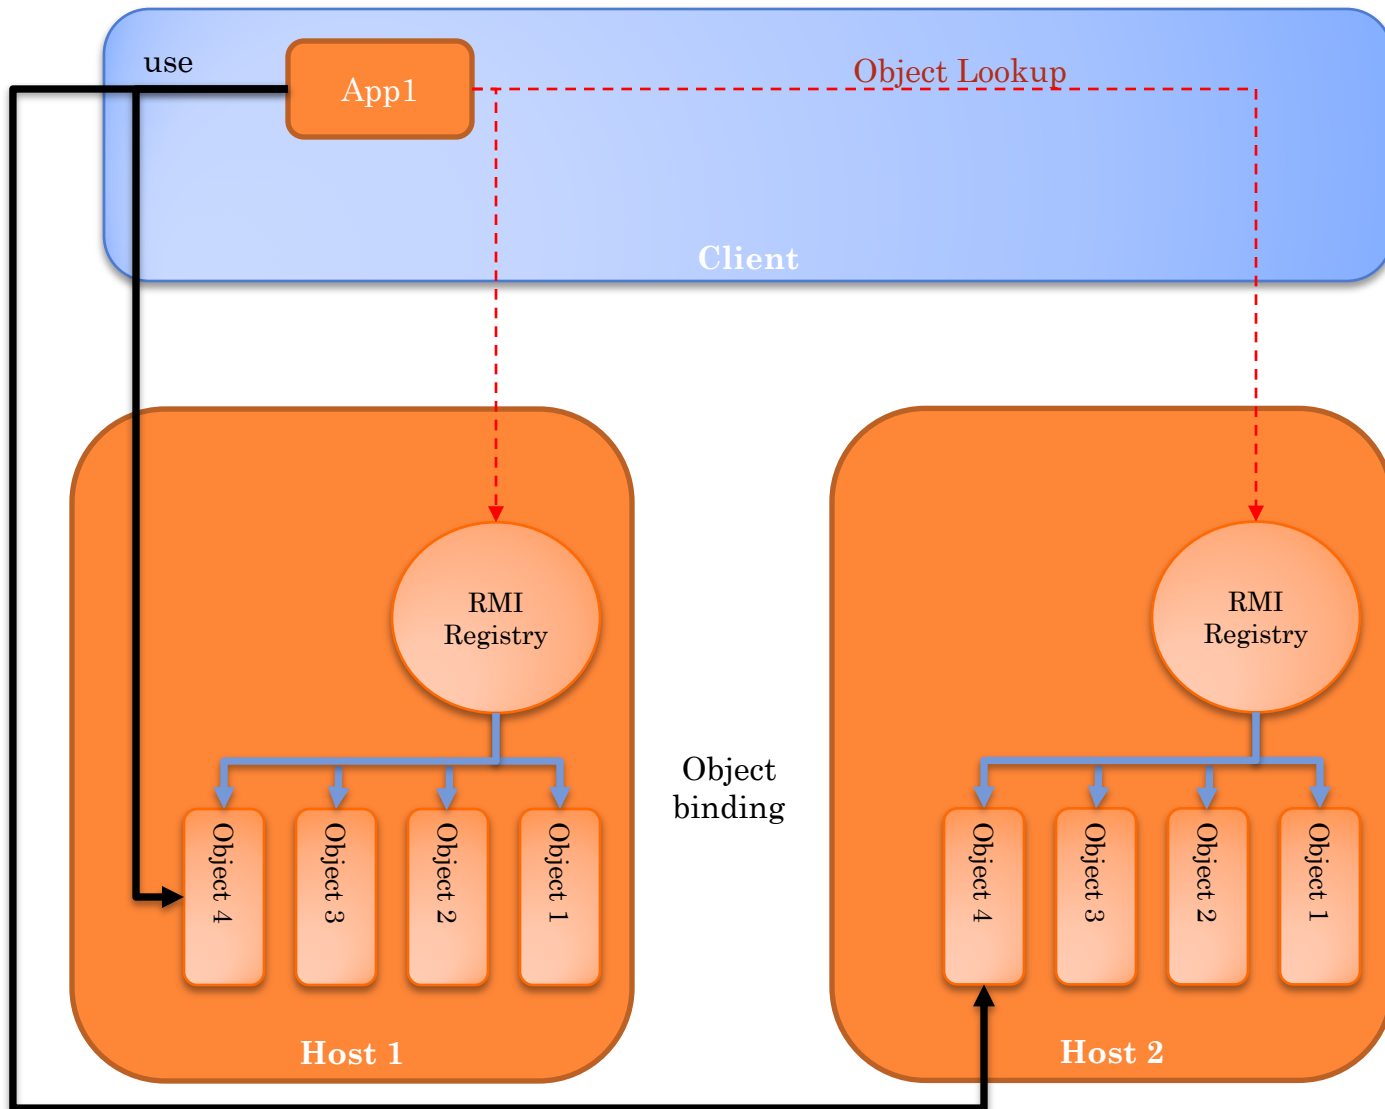
- **Key Design Principle:** Separate the implementation from the public interface.
 - Use Java interfaces to define the public interface of the remote object.
 - Implement the interface as a Java class.
- Export the implementation as a remote object.
 - Generates a **skeleton** to act as the remote object server.
 - Skeleton registers object with an RMI server (JRMP or IIOP)
 - Uses a **stub** (proxy) to act as the remote client.
 - Stub created using `java.lang.reflect.Proxy` class
 - Stub is a client that interacts with the skeleton.
- A global registry (naming service) is used to resolve object identifiers to remote object addresses (URLs).
 - Stub is stored in the registry and passed to the client for use



JAVA RMI ARCHITECTURE



JAVA RMI ARCHITECTURE





EXAMPLE: RMI CALCULATOR

RMI CALCULATOR

- Basic Calculator Service:
 - Declare a Remote Object Interface:

```
public interface Calculator extends java.rmi.Remote {  
    public long add(long a, long b) throws java.rmi.RemoteException;  
    public long sub(long a, long b) throws java.rmi.RemoteException;  
    public long mul(long a, long b) throws java.rmi.RemoteException;  
    public long div(long a, long b) throws java.rmi.RemoteException;  
}
```

- Implement the Remote Object:

```
public class CalculatorImpl implements Calculator {  
    public long add(long a, long b) { return a + b; }  
    public long sub(long a, long b) { return a - b; }  
    public long mul(long a, long b) { return a * b; }  
    public long div(long a, long b) { return a / b; }  
}
```



RMI CALCULATOR

○ Basic Calculator Service:

- Deploy the Remote Object:

```
public class CalculatorServer {  
    public static void main(String args[]) {  
        try {  
            // Create the RMI Registry  
            Registry registry = LocateRegistry.createRegistry(1099);  
  
            // Create the Remote Object  
            Calculator c = (Calculator)  
                UnicastRemoteObject.exportObject(new CalculatorImpl(), 0);  
  
            // Register the object with the RMI Registry  
            registry.bind("CalculatorService", c);  
        } catch (Exception e) {  
            System.out.println("Trouble: " + e);  
        }  
    }  
}
```




RMI CALCULATOR

○ Basic Calculator Service:

- Interact with Remote Object:

```
public class CalculatorClient {  
    public static void main(String[] args) {  
        try {  
            // Get a reference to the RMI Registry  
            Registry registry = LocateRegistry.getRegistry();  
  
            // Find the distributed object (stub created here)  
            Calculator c = (Calculator) registry.lookup("CalculatorService");  
  
            // Do stuff!!!!  
            System.out.println(c.sub(4, 3));  
            System.out.println(c.add(4, 5));  
            System.out.println(c.mul(3, 6));  
            System.out.println(c.div(9, 3));  
        } catch (Exception e) {  
            e.printStackTrace();  
        }  
    }  
}
```



RMI DECONSTRUCTED

- The RMI Registry is a central server that stores information on the remote objects that are available.
- Creating the RMI registry:
 - You can run the RMI registry separately by executing the `rmiregistry` program in `bin` folder of the JDK.
 - Alternatively, you can create it programmatically via the following statement:

```
Registry registry = LocateRegistry.createRegistry(1099);
```
- Getting the RMI registry:
 - Once the registry has been created, clients can access the registry via the following statement:

```
Registry registry = LocateRegistry.getRegistry();
```
 - Alternatively:

```
Registry registry = LocateRegistry.getRegistry("localhost",1099);
```



RMI DECONSTRUCTED

○ Creating a Remote Service:

- To do this, you simply create an instance of the Remote Service and pass it to the `exportObject(..)` method of the `UnicastRemoteObject` class.

```
Calculator c = (Calculator)  
    UnicastRemoteObject.exportObject(new CalculatorImpl(), 0);
```

○ Advertising a service:

- This makes the service available over RMI.
- To make a service available, you **bind** it to a name:

```
registry.bind("CalculatorService", c);
```

○ Finding a service:

- You need to search for a service with the same name as declared on the server:

```
Calculator c = registry.lookup("CalculatorService");
```



FULL REGISTRY API

- `void bind(String name, Remote obj)`
 - Binds the specified name to a remote object.
- `String[] list(String name)`
 - Returns an array of the names bound in the registry.
- `Remote lookup(String name)`
 - Returns a reference, a stub, for the remote object associated with the specified name.
- `void rebind(String name, Remote obj)`
 - Rebinds the specified name to a new remote object.
- `void unbind(String name)`
 - Destroys the binding for the specified name that is associated with a remote object.



RMI LIMITATIONS

- Only supports Java-to-Java interaction
- Lookup service limited to name.
 - What if you don't know the name?
 - What if you are looking for a type of service?
 - E.g. a weather service, a payment service, ...
- Interaction is strongly coupled:
 - You decide on the object (service) you want to interact with.
 - Client blocks until the response is returned.
 - What happens if the service fails?
 - What if the service becomes overloaded?
- Remember: Distributed Systems are expected to handle failure and should be designed to scale with demand...



The left side of the slide features a series of vertical stripes in shades of brown, tan, and grey. Overlaid on these stripes are several orange circles of varying sizes, arranged in a cluster that tapers towards the bottom.

NOTE ABOUT THE PROXY...

HOW DOES THE PROXY WORK?

- RMI works because the stub and skeleton can be automatically generated.
 - This is achieved by using some magic from the **Reflection API**.
 - This API provides a set of classes for representing Java classes, Interfaces, Methods, Annotations, ...
 - The most common use of the API is to create an instance of a class based on a string representation of its canonical name:

```
Object obj = Class.forName("java.lang.Object").newInstance();
```

- One very interesting class in this API is the `java.lang.reflect.Proxy` class.
- This class can be used to dynamically create an object that seems to be an instance of an explicitly specified set of Java interfaces – **even if no actual instance exists**.



HOW DOES THE PROXY WORK?

```
public class DebugProxy implements java.lang.reflect.InvocationHandler {  
    private Class cls;  
    public static Object newInstance(Class cls) {  
        return java.lang.reflect.Proxy.newProxyInstance(  
            cls.getClassLoader(), new Class[] {cls}, new DebugProxy(cls));  
    }  
    private DebugProxy(Class cls) { this.cls = cls; }  
    public Object invoke(Object proxy, Method m, Object[] args) throws Throwable {  
        System.out.println("class " + cls.getCanonicalName());  
        System.out.println("method " + m.getName());  
        System.out.println("args " + Arrays.asList(args));  
        return 01;  
    }  
}  
  
public class Main2 {  
    public static void main(String[] args) throws RemoteException {  
        Calculator calculator = (Calculator) DebugProxy.newInstance(Calculator.class);  
        calculator.add(50, 50);  
    }  
}
```



A PRACTICAL EXAMPLE...

```
public class DebugProxy implements java.lang.reflect.InvocationHandler {  
    private Object obj;  
    public static Object newInstance(Object obj) {  
        return java.lang.reflect.Proxy.newProxyInstance(obj.getClass().getClassLoader(),  
            obj.getClass().getInterfaces(), new DebugProxy(obj));  
    }  
    private DebugProxy(Object obj) { this.obj = obj; }  
    public Object invoke(Object proxy, Method m, Object[] args) throws Throwable {  
        Object result;  
        try {  
            System.out.println("before method " + m.getName());  
            result = m.invoke(obj, args);  
        } catch (InvocationTargetException e) {  
            throw e.getTargetException();  
        } catch (Exception e) {  
            throw new RuntimeException("unexpected invocation exception: " + e.getMessage());  
        } finally {  
            System.out.println("after method " + m.getName());  
        }  
        return result;  
    }  
}
```



A PRACTICAL EXAMPLE...

```
import java.rmi.RemoteException;

public class Main {
    public static void main(String[] args) throws RemoteException {
        Calculator calculator =
            (Calculator) DebugProxy.newInstance(new CalculatorImpl());
        System.out.println("Result: "+calculator.add(50, 50));
    }
}
```

----- EXPECTED OUTPUT -----

before method add

after method add

Result: 100



A decorative graphic on the left side of the slide. It features a series of vertical stripes in shades of brown, tan, and grey. Overlaid on these stripes are several orange circles of varying sizes, arranged in a cluster that tapers downwards.

APACHE RIVER (AKA JINI)

BRIEF HISTORY

- Developed by Sun Microsystems using Java
 - First Released in 1998
 - Ownership transferred to Apache in 2012
 - Still active development - current release 3.0.0
- One of the first attempts to develop a decentralised infrastructure for Object-Oriented distributed systems.
 - Uses a custom protocol for remote method invocation (JERI).
 - Based around triplet of concepts: **Registry**, **Service**, **Client**.
 - Includes a HTTP server to support on-demand distribution of compiled code (**Class Server**).
 - Packaged with services for: locking, distributed events, leasing, transactions, etc.
 - Includes a **tuple space** implementation called **JavaSpaces**.
 - Highly configurable architecture



LOOKUP SERVICE (LUS)

- Centralised service that consists of a set of **service items**:
 - **Unique identifier** is assigned to the service
 - **Service Type**: typically the fully qualified class name of the service.
 - **Associated set of attributes**: key-value pairs that define the properties of the service
 - **Client code**: downloadable interface to the service (e.g. RMI stub).
- Key Functionality:
 - Registration performed when a service is deployed.
 - Location of services can be based on id, type, attributes or some combination of the three.



WORKING A SERVICE

- Each service consists of:
 - An implementation (in Java):
 - Implements `net.jini.lookup.ServiceIDListener`
 - Service creation and deployment logic implemented in constructor based on service configuration.
 - A service configuration file (implemented as named collections of key-value pairs)
 - Declares how the client is created
 - Sets necessary permissions for security
 - Declares static attributes (for the LUS)
 - Defines how the service can be discovered
- Clients use the LUS to locate a service, download the client code, and invoke the remote method.



JAVASPACES

- Tuple Spaces are a shared memory model for distributed systems.
 - Based on the Blackboard concept from Distributed AI
 - Data is shared between nodes in the form of **tuples** – a series of typed fields.
 - Nodes are able to read/write tuples from/to a shared space independent of their physical location.
 - When reading, tuples are selected via pattern matching – you specify some of the data (a tuple template) and the tuple space matches it to a tuple.
- In the JavaSpaces framework:
 - Objects are used as tuples (type matters here)
 - Templates are also objects, where **null** means **any value**.



PAPER ALLOCATIONS

- A – HADOOP
- B – SYNAPSE
- C – CASSANDRA
- D – CHUBBY
- E – ZOOKEEPER
- F – DYNAMO
- G – SCRIBE
- H – CEPH
- I – BIGTABLE
- J – MEGASTORE
- K – VOLDEMORT
- L – SPANNER
- M – PERCOLATOR
- N – ACOP
- O – AMBRY
- P – ESPRESSO
- Q – RAFT
- R – ACTORDBS
- S – SPARK
- T – KAFKA

