



COURSE INTRODUCTION

COMP 30220: Distributed Systems

Lecturer: Rem Collier

Email: rem.collier@ucd.ie

WHAT IS THIS COURSE ABOUT?

- This course will explore some of the core practical concepts that underpin **distributed systems**.
 - Basic Networking (Sockets)
 - Distributed Objects (RMI)
 - Web Services / SOAP & REST
 - Message-Oriented Middleware / Active MQ + JMS
- We will also try to touch on some “bigger” issues:
 - Enterprise Service Bus / Choreography & Orchestration
 - Microservices / Service Mashups
 - Reactive Systems / Actor-based Systems



COURSE ASSESSMENT

- This course is 100% assessment:
 - Paper Review 20%
*Read an assigned papers and discuss in a reading group.
Submit individual reports.*
 - Group Presentation 20%
Reading Groups come together to present assigned papers.
 - Practical Work 20%
Preparatory programming work.
 - Group Project 40%
Large Student-Driven Project that uses techniques from module.
- Labs will start in week 2, papers and reading groups will be assigned in week 3.



COURSE SCHEDULE

Week	Lectures	Labs	Notes
1	Introduction	n/a	
2	Network Programming	Sockets	
3	Distributed Objects	Java RMI	Paper Assignment
4	Web Services / SOAP	Jax-WS	
5	Web Services / REST	Restlets	
6	Message-Oriented Middleware	READING GROUPS	Report Submission
7	Actors	JMS + Apache MQ	Proposal Submission
8	4x PRESENTATIONS	4x PRESENTATIONS	
9	4x PRESENTATIONS	4x PRESENTATIONS	
10		PROJECT SUPPORT	
11		PROJECT SUPPORT	
12		PROJECT SUPPORT	
13	--READING WEEK--		Project Submission



EXAMPLE PROJECTS


- **Spy Polite:** proxy system for Facebook to allow secure sharing of users credentials / accessing photos of friends of friends.
- **Rotten Tomatoes:** Crowdsourced YouTube service based on REST + GraphQL.
- **Spaceship Tamogatchi:** REST-based tamogatchi space game.
- **Games4gamers:** JMS-based game server & client
- **Distributed Betting Service:** REST-based mashup of betting sites to provide single point of reference.
- **Billy Rubin & the Jets:** JMS-based Distributed File Service



ACCESS TO COURSE MATERIAL

- <http://csmoodle.ucd.ie/>
- Login using UCD credentials
- COMP30220 Distributed Systems 2018-19
- Enrolment Key: “ds18”





INTRODUCTION

A SHORT HISTORY OF “DISTRIBUTED” COMPUTING

○ Humble Beginnings:


- Bankomat ATM (1968)
- ARPANET (1969) - File Transfer / Email
- **Remote Procedure Calls (1970s)**
- The Internet (1982) - TCP/IP
- Unix NFS (1980-) / Andrew File System (1982-)
- The World Wide Web (1991) - URI's, HTML, HTTP
- **CORBA: Common Object Request Broker Architecture (1991)**

○ The Rise of the Internet:


- Web Browsers: Mosaic (1993), Netscape (1994), IE (1996)
- E-commerce: Amazon (1994), ebay (1995), PayPal (1998)
- Search Engines: Excite (1993), Yahoo (1994), Google (1998)
- **IIOP (1996), Java RMI (1997), XML-RPC (1998)**
- P2P: Napster (1999), Gnutella (2000), BitTorrent (2001)

○ Web 2.0:

- Social Media: LinkedIn (2002), MySpace (2003), Facebook (2004), Twitter (2008), WhatsApp (2009), Snapchat (2011)
- Mobile Apps / Gaming (2008-)
- **SOAP (2000), REST (2000), JSON (2001), JMS (2002), JSON-RPC (2005)**



GROWTH OF THE INTERNET



WHY DISTRIBUTED SYSTEMS?

○ **Cost/Performance/Scalability.**

- Better price/performance as long as commodity hardware is used for the component computers.
- Resources such as processing and storage capacity can be increased incrementally.
- Scale Up vs Scale Out

○ **Reliability/Fault Tolerance.**

- Redundancy of functionality / replication of data.

○ **Inherent distribution.**

- Some applications, such as email and the Web (where users are spread out over the whole world), are naturally distributed.
- This includes cases where users are geographically dispersed as well as when single resources (e.g., printers, data) need to be shared.




HOWEVER...

- **New component: network.** Networks are needed to connect independent nodes and are subject to performance limitations. Besides these limitations, networks also constitute new potential points of failure.
- **Security.** Because a distributed system consists of multiple components there are more elements that can be compromised and must, therefore, be secured. This makes it easier to compromise distributed systems.
- **Software complexity.** As will become clear throughout this course - distributed software is more complex and harder to develop than conventional software; this makes it more expensive to develop and there is a greater chance of introducing errors.



EXAMPLE

- A bank asks you to program their new ATM software
 - Central bank computer (server) stores account information
 - Remote ATMs authenticate customers and deliver money



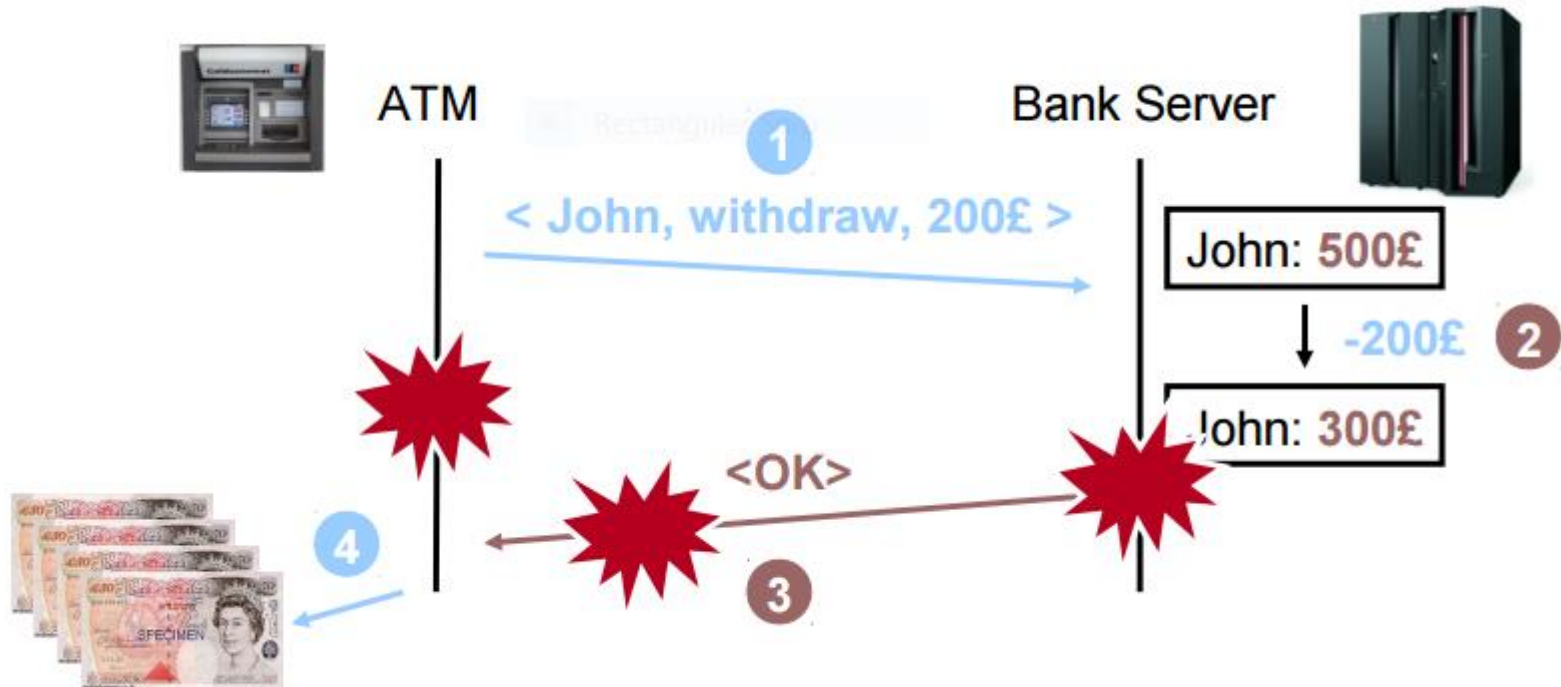
EXAMPLE: NAÏVE SOLUTION

- ATM: (ignoring authentication and security issues)
 1. Ask customer how much money s/he wants
 2. Send message with **<customerId, withdraw, amount>**, to bank server
 3. Wait for bank server answer: **<OK>** or **<refused>**
 4. If **<OK>** give money to customer, else display error message

- Central Server:
 1. Wait for messages from ATM:
 2. If enough money withdraw money, send **<OK>**, else send **<refused>**



EXAMPLE: NAÏVE SOLUTION



○ But ...

- What if the bank server crashes just after 2 and before 3?
- What if the message gets lost? Takes days to arrive?
- What if the ATM crashes after 1, but before 4?

TYPICAL FALSE ASSUMPTIONS

- The network is reliable
- The network is secure
- Everything is homogeneous
- The topology does not change
- Latency is zero
- Bandwidth is infinite
- Transport cost is zero
- There is one administrator



THE CAP (BREWERS) THEOREM

- It is impossible for a distributed data store to satisfy all three of the following guarantees:
 - **Consistency:** Every read receives the most recent write or an error
 - **Availability:** Every request receives a non-error response – without guarantee that it contains the most recent write
 - **Partition tolerance:** The system continues to operate despite an arbitrary number of messages being dropped (or delayed) by the network between nodes.
- In short, designers of distributed data stores must choose between consistency and availability:
 - SQL databases: Consistency over Availability (ACID)
 - NoSQL databases: Availability over Consistency (BASE)



CONSISTENCY MODELS

- **Strict Consistency:** a write to a variable / data store must be seen **immediately** by all nodes.
 - In practice can only be achieved by centralised systems.
- **Sequential Consistency:** writes to variables / data stores must be seen in the **same order**.
 - Requires some form of global clock to ensure a consistent ordering
 - Can lead to poor performance in practice
 - Sometimes necessary – e.g. banking systems
- **Eventual Consistency:** reads eventually return the same value (there is no guarantee that it is the latest write).
 - Good for systems where sequential consistency is not required (messages)
 - For example – Twitter Tweets follow eventual consistency (your tweet eventually appears in all your followers feeds).
- **Key Issue:** How to know the order of the writes?



INSERT: GLOBAL TIME

- Timestamps provide an ultimate ordering, but...
 - Local clocks are not consistent – clock drift
 - Berkeley and Cristian's Algorithms can fix it on the intranet...
 - BUT –still have to deal with latency, failure, ...
- Universal Coordinated Time (UTC)
 - Global hierarchy of clock servers rooted in a set of synchronised Cesium-133 clocks hosted by accepted authorities.
 - Setting on every computer to synchronise with a Network Time Protocol (NTP) signal that is broadcast on most networks.
- Get Logical
 - Stop trying to get the real time, use logical time instead
 - Lamport Clocks – Synchronisation based on **Happens Before**
 - Vector Clocks – Time is a set of counters (1 per machine)




AVAILABILITY

- Availability in the presence of network failure requires replication of data.
 - Each replicas state must be maintained when any update occurs.
- Sequential consistency requires that all updates must be applied in the same order and become visible at the same time.
 - If you have 2 replicas of a piece of data, D, on machines A and B, and D is updated on machine A, then the copy of D on machine B **must be updated before** the update can be read on machine A.
- Eventual consistency requires that all updates eventually be applied (in the same order).
 - In our example above, the A can allow reads of the updated D **before the copy of D has been updated on B.**



FINALLY: REMOTE PROCEDURE CALLS



- Form of Socket-Based **Inter-Process Communication** that mimics a procedure call.
 - Interaction is **synchronous**, based on a **request-response protocol**.
 - The processes are known as the **client** and the **server**.
 - A **client stub** constructs the request and parses the response.
 - A **server stub** parses the request; invokes the procedure and constructs the response.
 - The conversion of parameters/return values to/from some intermediary format is known as **marshalling/unmarshalling**.
 - The intermediary format is known as an Interface Description Language (IDL).