



NETWORK PROGRAMMING: A JAVA PERSPECTIVE

COMP 30220: Distributed Systems

Lecturer: Rem Collier

Email: rem.collier@ucd.ie

TCP vs. UDP DATAGRAM

- TCP is designed for reliable communications and ensures that lost or damaged data is resent
 - if packets arrive out of order TCP puts them back into the correct order
 - if the receiver is being flooded by data arriving too quickly TCP slows down the transmission to prevent packets being lost
- The User Datagram Protocol (UDP) is not reliable
- UDP datagrams add very little to the IP datagrams they sit on top of:
 - The header adds 8 bytes for storing source and destination ports
 - The number of bytes of data following the IP header
 - And an optional checksum...



UDP DATAGRAM

- UDP doesn't provide any of the checks that TCP does
 - These checks are useful but slow down TCP
 - UDP is **lean** and **fast** but **error prone**
- If you require reliable transmission (e.g. FTP) you wouldn't normally use UDP.
- If you want fast transmission and don't care about a lost or swapped packet here or there then you would consider UDP (e.g. streaming real-time audio or video)
- You should also consider how reliable your network is...




EXAMPLE: UDPSERVER

```
class UDPServer {  
    public static void main(String args[]) throws Exception {  
        DatagramSocket serverSocket = new DatagramSocket(9876);  
  
        while (true) {  
            // Wait for a packet  
            DatagramPacket receivePacket = new DatagramPacket(new byte[1024], 1024);  
            serverSocket.receive(receivePacket); // Blocking Method  
  
            // Process the received data  
            String sentence = new String(receivePacket.getData());  
            System.out.println("RECEIVED: " + sentence);  
  
            // Generate response and transmit  
            byte[] sendData = sentence.toUpperCase().getBytes();  
            DatagramPacket sendPacket = new DatagramPacket(sendData, sendData.length,  
                                                            receivePacket.getAddress(), receivePacket.getPort());  
            serverSocket.send(sendPacket);  
        }  
    }  
}
```



EXAMPLE: UDPCIENT

```
class UDPClient {  
    public static void main(String args[]) throws Exception {  
        DatagramSocket clientSocket = new DatagramSocket();  
  
        // Read User Input (and convert to byte array for transmission)  
        BufferedReader inFromUser = new BufferedReader(new InputStreamReader(System.in));  
        byte[] sendData = inFromUser.readLine().getBytes();  
  
        // Send the packet  
        DatagramPacket sendPacket =  
            new DatagramPacket(sendData, sendData.length, InetAddress.getByName("localhost"), 9876);  
        clientSocket.send(sendPacket);  
  
        // Wait for a response...  
        DatagramPacket receivePacket = new DatagramPacket(new byte[1024], 1024);  
        clientSocket.receive(receivePacket);  
  
        // Display the output  
        System.out.println("FROM SERVER:" + new String(receivePacket.getData()));  
        clientSocket.close();  
    }  
}
```



UDP DECONSTRUCTED

○ Creating a socket:

- Server:

```
DatagramSocket socket = new DatagramSocket(<port>);
```

- Client:

```
DatagramSocket socket = new DatagramSocket();
```

○ Sending data:

```
byte[] data = new byte[1024];
```

```
DatagramPacket packet =
```

```
    new DatagramPacket(data, data.length, <host>, <port>);
```

```
socket.send(packet);
```

○ Receiving data:

```
byte[] data = new byte[1024];
```

```
DatagramPacket packet = new DatagramPacket(data, data.length);
```

```
socket.receive(packet);
```



UDP DECONSTRUCTED

- Closing a socket:

```
socket.close();
```

- UDP Packet Operations:

- Getting the source address:

```
InetAddress IPAddress = packet.getAddress();
```

- Getting the source port:

```
int port = packet.getPort();
```

- Getting the data:

```
String modifiedSentence = new String(packet.getData());
```



MULTICAST

- Previous two examples were point-to-point (unicast)
- Some tasks require 1-many communications
 - E.g. streaming radio and television
- But broadcast is too indiscriminate
 - Especially in switched networks
- Multicast fills the gap
 - it is broader than point-to-point unicast communications but narrower than broadcast
- Clients express interest in multicast data by joining a particular multicast group



MULTICAST SOCKETS

- Multicast is implemented by **multicast sockets**
 - Multicast sockets are similar to UDP sockets with the exception that they express interest in a multicast group rather than a specific port.
 - Multicast addresses are IP addresses in the range 224.0.0.0 to 239.255.255.255.
 - This range was chosen because all addresses have the binary digits 1110 as their first four bits
 - Collectively they are known as **class D addresses**
 - Like any IP address multicast addresses can have a hostname
 - for example 224.0.1.1 (the network time protocol service) is ntp.mcast.net
- Multicast Applications:
 - Audio and Video
 - Multiplayer (network) games
 - distributed file systems
 - Conferencing
 - database replication
 - Multicast DNS (mDNS / Bonjour)
- Multicast is typically seen as a **LAN-level service** due to bandwidth concerns.



MULTICAST ADDRESSES

- Network nodes can join together and form a multicast group all represented by a single multicast address.
- Multicast addresses are IP addresses in the range 224.0.0.0 to 239.255.255.255
 - All addresses in this range have the binary digits 1110 as their first four bits and are collectively known as class D addresses.
 - Like any IP address multicast addresses can have a hostname
 - For example 224.0.1.1 (the Network Time Protocol service) is ntp.mcast.net
 - Some IP addresses are reserved by the Internet Assigned Numbers Authority (IANA)
 - You can create your own transient **multicast group** by picking a random address between 225.0.0.0 and 238.255.255.255
 - There is no multicast address that sends data to all hosts on the Internet!



SENDING MULTICAST DATA

- When a host wants to send to a multicast group it puts the data into a UDP Datagram packet
 - This is done for speed over reliability.
 - Connection-orientated TCP would be a bad choice for the multicast protocols because of ACK explosion.
 - Reliable Multicast Services do exist.
- From the programmers perspective it looks pretty much the same as UDP communications
 - However you have to worry about the time-to-live (TTL) value
 - The TTL says how many router hops a packet may take before being discarded
 - The TTL in multicast geographically limits the range of the packets.




MULTICAST RECEIVER IN JAVA

```
public class MCastReceiver {  
    public static void main(String[] args) {  
        byte[] buffer = new byte[1024];  
        DatagramPacket gift = new DatagramPacket(buffer, buffer.length);  
        try {  
            InetAddress address = InetAddress.getByName("225.0.0.1");  
            MulticastSocket ms = new MulticastSocket(5000);  
            ms.joinGroup(address);  
            while (true) {  
                ms.receive(gift);  
                String present =  
                    new String(gift.getData(), 0, gift.getLength());  
                System.out.println(present);  
            }  
        } catch (Exception e) {}  
    }  
}
```



MULTICAST SENDER IN JAVA

```
public class MCastSender {  
    public static void main(String[] args) {  
        byte[] giftData = "Santa has got you a new bike".getBytes();  
        try {  
            // Join a multicast group  
            MulticastSocket ms = new MulticastSocket();  
            InetAddress address = InetAddress.getByName("225.0.0.1");  
            ms.joinGroup(address);  
            ms.setTimeToLive(1);  
  
            // Send some data  
            ms.send(new DatagramPacket(giftData, giftData.length, address, 5000));  
  
            // Leave the group  
            ms.leaveGroup(address);  
            ms.close();  
        } catch (Exception e) {}  
    }  
}
```





DYNAMIC HEALTH/DISCOVERY SERVICE

- Each node joins a on a specified group and listens on a predefined port.
- Nodes periodically broadcast a “HELLO” message to the group.
 - The IP Addresses are used to identify the nodes AND that they are still available.
 - Failure to receive a “HELLO” in a given period of time can indicate that the node is no longer available.
- Variations on this theme can be used to offer enhanced services:
 - The hello message can include point-to-point contact info.
 - Bridges can be used to link multiple multicast networks.



A THIRD WAY: JAVA NIO

- Part of Java since v1.4 (2002)
- Socket based implementations are slow:
 - Threads are expensive (~2Mb per thread)
 - Most CPU time is spend blocking (waiting to read)
- Java NIO (New IO) is an alternative IO API for Java.
 - Decouples data processing from connection handling
 - Read/Writes occur in parallel with data processing
 - Designed to scale
- Based on Channels, Buffers and Selectors:
 - Data is read from a buffer into a channel
 - Data is written from a channel into a buffer
 - Selectors monitor channels and handle “events”



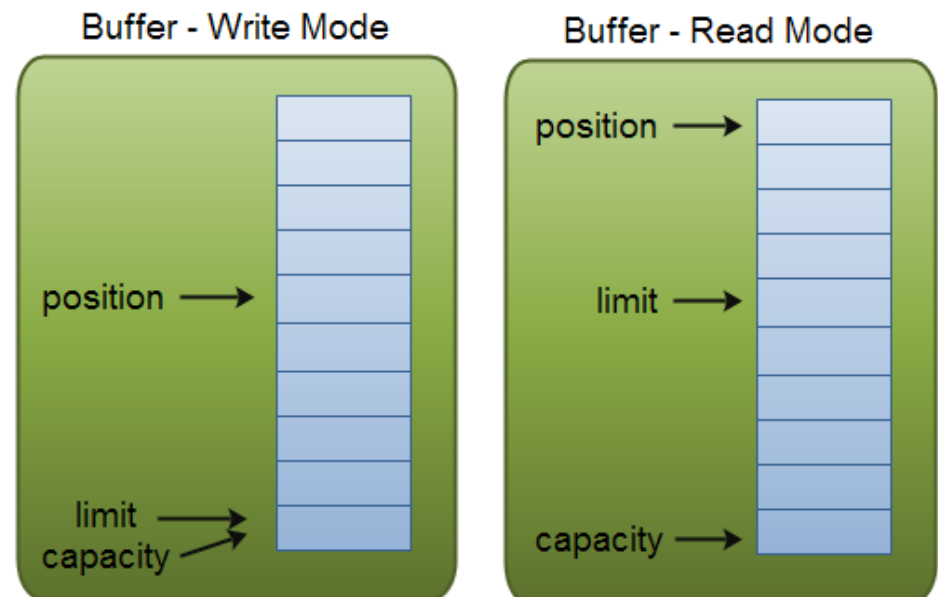
JAVA NIO: CHANNELS

- Channels are like streams:
 - Data can be written to a channel or read from a channel.
 - Some commonly used channels are:
 - `FileChannel`
 - `DatagramChannel`
 - `SocketChannel`
 - `ServerSocketChannel`
- Channels are sources (or destinations) for the data being processed.
 - Like a stream you can read from or write to a channel.
- Channels generate events that can be monitored by a selector:
 - Connect, Accept, Read, Write



JAVA NIO: BUFFER DESIGN

- Buffers can be in either read or write mode.
- Buffers are typed (all Java primitives supported)
 - E.g. `ByteBuffer`, `LongBuffer`
- Three key pieces of information:
 - **Capacity**: the maximum number of values that can be stored in the buffer.
 - **Position**: the location of the next read / write in the buffer.
 - **Limit**: how much you can actually read/write for the buffer.



JAVA NIO: BASIC CONCEPT

- NIO Buffers have a mode:
 - The can be set to **read** or **write** mode.
 - The program defines when the buffer changes mode.
- Using a **Buffer** involves 5 steps:
 - Create a buffer (in write mode)
 - Write data to a buffer
 - Flip the buffer (to read mode)
 - Read data from the buffer
 - Clear / Compact the buffer and return to write mode
- Other functions of a buffer include:
 - Rewinding, marking and resetting, ...



JAVA NIO: READING FROM A FILE

```
RandomAccessFile aFile =
    new RandomAccessFile("data/nio-data.txt", "rw");

FileChannel inChannel = aFile.getChannel(); // Create a FileChannel
ByteBuffer buf = ByteBuffer.allocate(48); // Buffer with 48 bytes (Write mode)

// Read from FileChannel to Buffer
int bytesRead = inChannel.read(buf);

// Check if any bytes were read...
while (bytesRead != -1) {
    System.out.println("Read " + bytesRead);

    // Flip Buffer to Read mode
    buf.flip();

    // Read from Buffer and write to Console (1 byte at a time)
    while(buf.hasRemaining()){
        System.out.print((char) buf.get());
    }

    // Empty Buffer and return to Write mode
    buf.clear();

    // Read more from FileChannel to Buffer
    bytesRead = inChannel.read(buf);
}

aFile.close();
```



JAVA NIO: WRITING TO A FILE

```
String newData = "New String to write to file..." + System.currentTimeMillis();
```

```
// Create a buffer
```

```
ByteBuffer buf = ByteBuffer.allocate(48);
```

```
// Clear it
```

```
buf.clear();
```

```
// Write data to the buffer
```

```
buf.put(newData.getBytes());
```

```
// Flip it to read mode
```

```
buf.flip();
```

```
// read data from the buffer..
```

```
while(buf.hasRemaining()) {
```

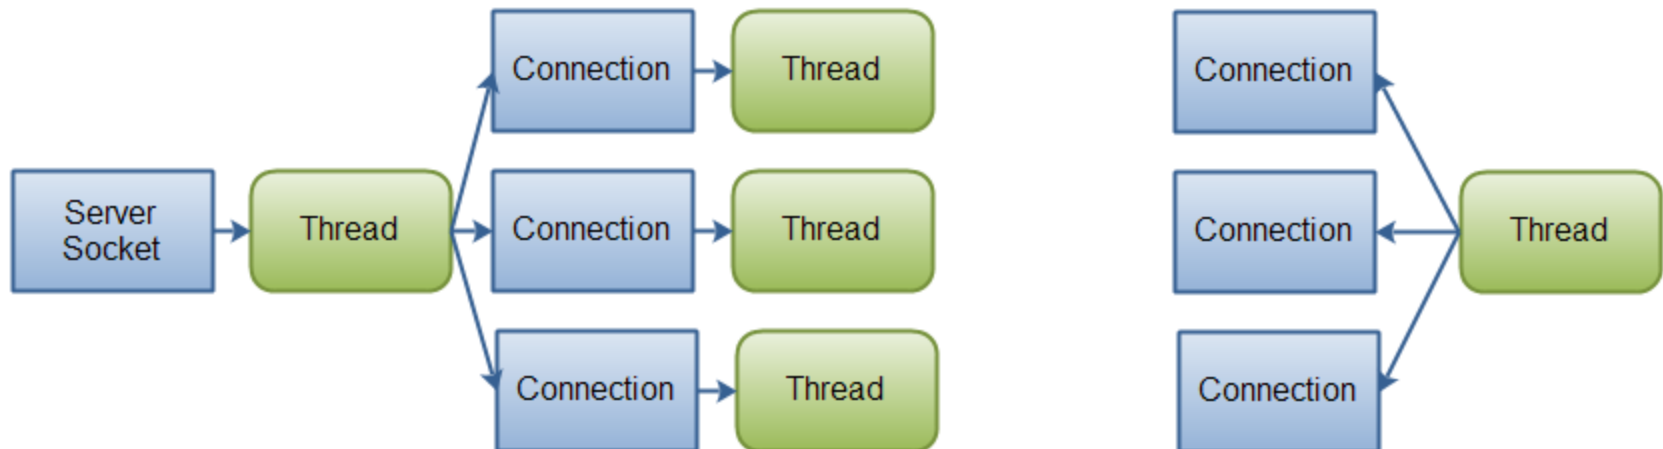
```
    channel.write(buf);
```

```
}
```



JAVA NIO: SELECTORS

- Component that manages multiple channels:
 - Channels register with a selector identifying which events the selector should handle.
 - The selector polls the registered channels whenever the `select()` method is invoked.
 - The selector returns a list of the “ready” channels that can then be processed.



JAVA NIO IN ACTION

- Java NIO Support:
 - Native Java support
 - Basic Library for components
 - <http://tutorials.jenkov.com/java-nio/>
- Netty:
 - Alternate Java implementation
 - Includes: HTTP, SMTP, FTP support out of the box
 - <https://netty.io>



OTHER TECHNOLOGIES

- Socket based communication is low layer and requires conversion between language specific data types and byte representations/strings for transmission.
- There are alternatives that attempt to alleviate the strain of coding bespoke solutions:
 - Common Object Request Broker Architecture (CORBA)
 - Language-neutral middleware
 - Java Remote Method Invocation (RMI)
 - Java-based middleware solution
 - Web Services
 - XML, SOAP-based architecture, JSON, REST
 - Java Messaging Service (JMS)
 - Message Oriented Middleware

